

Name: Sadneya Sadanand Samant

Roll no:46 Class:D15C

EXPERIMENT NO: 1

Aim: Introduction to Data science and Data preparation using Pandas steps.

Theory:

Data preparation is a crucial step in data science, involving cleaning and transforming raw data into an analyzable format. Using Pandas, we can perform operations such as handling missing values, encoding categorical data, and scaling numerical features. Proper preprocessing ensures the dataset is reliable for analysis and modeling by addressing inconsistencies, missing data, and outliers.

Problem Statement:

The Vehicle Safety Recall dataset, provided by NHTSA, contains 15 columns detailing various aspects of recall events, such as manufacturers, affected components, and corrective actions. This analysis focuses on:

- **Manufacturer Trends:** Identifying manufacturers prone to frequent recalls or specific defects.
- **Impact Analysis:** Understanding recall types affecting the largest populations and assessing average completion rates.
- **Temporal Patterns:** Detecting trends in recalls over time and seasonal spikes.
- **Safety Implications:** Investigating critical safety advisories like "Do Not Drive" or "Park Outside" and their resolution rates.

By cleaning the dataset and applying data preprocessing steps, the goal is to enhance its quality and draw actionable insights for stakeholders.

Dataset Overview:

The dataset provides detailed information about vehicle safety recalls managed by the National Highway Traffic Safety Administration (NHTSA). It contains 15 columns, each capturing specific aspects of recall events. Below is a breakdown of the columns and their relevance:

1. **Report Received Date:** Date the recall was officially reported.
2. **NHTSA ID:** A unique identifier for each recall event.
3. **Recall Link:** A hyperlink to the recall details on the NHTSA website.
4. **Manufacturer:** Name of the vehicle or product manufacturer responsible for the recall.

- 5. Subject:** Brief description of the recall issue.
- 6. Component:** The affected part of the vehicle/product (e.g., "POWER TRAIN").
- 7. Mfr Campaign Number:** Manufacturer's internal reference for the recall.
- 8. Recall Type:** Type of product involved (e.g., vehicle, tire, or car seat).
- 9. Potentially Affected:** Number of units potentially impacted by the recall.
- 10. Recall Description:** Detailed explanation of the defect or issue.
- 11. Consequence Summary:** Description of the risks or consequences associated with the defect.
- 12. Corrective Action:** Steps taken to address the defect.
- 13. Park Outside Advisory:** Indicates whether there's an advisory to park outside for safety.
- 14. Do Not Drive Advisory:** Indicates whether there's an advisory not to drive the affected vehicle.
- 15. Completion Rate %:** Percentage of affected vehicles repaired or addressed.

Steps:

1. Loading The Dataset

```

✓ 2s [1] import pandas as pd
✓ 0s [2] df = pd.read_csv('recalls.csv')

```

2. Description of the dataset

a. Information about dataset

```

▶ df.info()

→ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 28671 entries, 0 to 28670
Data columns (total 15 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Report Received Date    28671 non-null   object 
 1   NHTSA ID            28671 non-null   object 
 2   Recall Link          28671 non-null   object 
 3   Manufacturer        28671 non-null   object 
 4   Subject             28671 non-null   object 
 5   Component           28671 non-null   object 
 6   Mfr Campaign Number  28624 non-null   object 
 7   Recall Type          28671 non-null   object 
 8   Potentially Affected 28630 non-null   float64
 9   Recall Description    26270 non-null   object 
 10  Consequence Summary  23783 non-null   object 
 11  Corrective Action     26283 non-null   object 
 12  Park Outside Advisory 28671 non-null   object 
 13  Do Not Drive Advisory 28671 non-null   object 
 14  Completion Rate % (Blank - Not Reported) 10007 non-null   float64
dtypes: float64(2), object(13)
memory usage: 3.3+ MB

```

b. Description of Dataset

```
# Get the dataset's shape and basic statistics
print(f"Dataset Shape: {df.shape}")
print(df.describe(include='all'))
```

```
Dataset Shape: (28671, 15)
   Report Received Date    NHTSA ID \
count          28671        28671
unique         10023        28671
top           10/17/2013  25E002000
freq            42             1
mean           NaN           NaN
std            NaN           NaN
min            NaN           NaN
25%           NaN           NaN
50%           NaN           NaN
75%           NaN           NaN
max            NaN           NaN
```

	Recall Link \
count	28671
unique	28671
top	Go to Recall (https://www.safercar.gov/recalls?nh...)
freq	1
mean	NaN
std	NaN
min	NaN
25%	25%
50%	50%
75%	75%
max	NaN

```
Mfr Campaign Number Recall Type Potentially Affected \
count          28624        28671      2.863000e+04
unique         11341             4           NaN
top           NR (Not Reported)  Vehicle           NaN
freq            16602        24940           NaN
mean           NaN           NaN       4.572011e+04
std            NaN           NaN       3.730381e+05
min            NaN           NaN      0.000000e+00
25%           NaN           NaN      9.900000e+01
50%           NaN           NaN      6.860000e+02
75%           NaN           NaN      6.385500e+03
max            NaN           NaN      3.200000e+07
```

	Recall Description \
count	26270
unique	25523
top	ON CERTAIN TRAILERS EQUIPPED WITH SEALCO SPRIN...
freq	28
mean	NaN
std	NaN
min	NaN
25%	NaN
50%	NaN
75%	NaN
max	NaN

```
Consequence Summary \
count            23783
unique          17015
top  RELEASE OF COOLANT UNDER CERTAIN CONDITIONS CO...
freq             128
mean           NaN
std            NaN
min            NaN
25%           NaN
50%           NaN
75%           NaN
max            NaN
```

	Corrective Action \
count	26283
unique	25579
top	DEALERS WILL EQUIP AIR SYSTEMS WITH A PRESSURE...
freq	18
mean	NaN
std	NaN
min	NaN
25%	NaN
50%	NaN
75%	NaN
max	NaN

```
Park Outside Advisory  Do Not Drive Advisory \
count            28671        28671
unique           2             2
top             No            No
freq            28601        28510
mean           NaN           NaN
std            NaN           NaN
min            NaN           NaN
25%           NaN           NaN
50%           NaN           NaN
75%           NaN           NaN
max            NaN           NaN
```

	Completion Rate % (Blank - Not Reported)
count	10007.000000
unique	NaN
top	NaN
freq	NaN
mean	67.874214
std	29.937993
min	0.000000
25%	48.350000
50%	76.390000
75%	93.765000
max	100.000000

3. Drop columns that aren't useful.

```
# Remove leading/trailing spaces from column names
df.columns = df.columns.str.strip()

# List of columns to drop
cols = ["Recall Link", "Mfr Campaign Number", "Park Outside Advisory", "Do Not Drive Advisory", "Completion Rate % (Blank - Not Reported)"]
```

```
# Drop the columns that are present in the DataFrame
df = df.drop(cols, axis=1)

# Display the updated DataFrame
print(df.head())
```

	Report Received Date	NHTSA ID	Manufacturer	Recall Description
0	01/14/2025	25E002000	GKN Automotive	0 GKN Automotive (GKN) is recalling certain repl...
1	01/13/2025	25E001000	N&B Mobility Solutions LLC	1 N&B Mobility Solutions LLC (Nivion) is recalli...
2	01/13/2025	25V005000	Forest River, Inc.	2 Forest River, Inc. (Forest River) is recalling...
3	01/13/2025	25V006000	Kia America, Inc.	3 Kia America, Inc. (Kia) is recalling certain 2...
4	01/13/2025	25V007000	Winnebago Industries, Inc.	4 Winnebago Industries, Inc. (Winnebago) is recal...
	Subject	Component	Consequence Summary	Corrective Action
0	Driveshaft Can Break	POWER TRAIN	0 A cracked or broken driveshaft can cause a los...	0 GKN will reimburse the cost of a replacement d...
1	Charger Adapter May Cause Arcing or Shock Risk	ELECTRICAL SYSTEM	1 Inadequate clearance between DC busbars may ca...	1 Nivion will replace the defective adapters, fr...
2	Cooktop Burner Tube May Crack and Cause Gas Leak	EQUIPMENT	2 A gas leak in the presence of an ignition sour...	2 Owners are advised not to use the cooktop until...
3	Loss of Headlights and Taillights/FMVSS 108	ELECTRICAL SYSTEM	3 A loss of headlights and taillights can reduce...	3 Dealers will update the BDC software, free of ...
4	Spare Tire Carrier May Detach	EQUIPMENT	4 A detached spare tire carrier can become a roa...	4 Dealers will inspect, replace, and correctly t...
	Recall Type	Potentially Affected		
0	Equipment	18.0		
1	Equipment	130.0		
2	Vehicle	396.0		
3	Vehicle	74469.0		
4	Vehicle	107.0		

Thus the columns now present in dataset are:

```
df.info()

→ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 28671 entries, 0 to 28670
Data columns (total 10 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Report Received Date    28671 non-null   object 
 1   NHTSA ID            28671 non-null   object 
 2   Manufacturer        28671 non-null   object 
 3   Subject             28671 non-null   object 
 4   Component           28671 non-null   object 
 5   Recall Type          28671 non-null   object 
 6   Potentially Affected 28630 non-null   float64 
 7   Recall Description   26270 non-null   object 
 8   Consequence Summary 23783 non-null   object 
 9   Corrective Action    26283 non-null   object 
dtypes: float64(1), object(9)
memory usage: 2.2+ MB
```

4. Take care of missing data.

a. Drop rows with maximum missing values.

```
print(f"Dataset Shape before Dropping Rows: {df.shape}")
# Drop rows with the highest number of missing values
threshold = len(df.columns) * 0.5 # Drop rows where over 50% of columns are missing
df = df.dropna(thresh=threshold)

print(f"Dataset Shape After Dropping Rows: {df.shape}")

→ Dataset Shape before Dropping Rows: (28671, 10)
Dataset Shape After Dropping Rows: (28671, 10)
```

```
▶ print(df.isnull().sum())

[1] Report Received Date      0
    NHTSA ID                  0
    Manufacturer               0
    Subject                   0
    Component                 0
    Recall Type                0
    Potentially Affected      41
    Recall Description          2401
    Consequence Summary        4888
    Corrective Action           2388
    dtype: int64
```

b. Handle Missing Data

Here above info says Potential Affected ,Recall Description ,Consequence Summary and corrective action contain some null values thus we need to handle missing data.

```
[12] # Fill missing numerical values with the median
     df['Potentially Affected'] = df['Potentially Affected'].fillna(df['Potentially Affected'].median())
     # Fill missing categorical values with a placeholder
     df['Recall Description'] = df['Recall Description'].fillna('Not Known')
     df['Consequence Summary'] = df['Consequence Summary'].fillna('Unknown')
     df['Corrective Action'] = df['Corrective Action'].fillna('Unknown')

     print(df.isnull().sum()) # Verify no missing values remain
```

```
[13] Report Received Date      0
    NHTSA ID                  0
    Manufacturer               0
    Subject                   0
    Component                 0
    Recall Type                0
    Potentially Affected      0
    Recall Description          0
    Consequence Summary        0
    Corrective Action           0
    dtype: int64
```

5. Create dummy variables

```
▶ # Convert categorical columns into dummy variables
     df = pd.get_dummies(df, columns=['Recall Type'], drop_first=True)

     print(df.head())
```

```

Report Received Date NHTSA ID Manufacturer \
0      01/14/2025 25E002000 GKN Automotive
1      01/13/2025 25E001000 N&B Mobility Solutions LLC
2      01/13/2025 25V005000 Forest River, Inc.
3      01/13/2025 25V006000 Kia America, Inc.
4      01/13/2025 25V007000 Winnebago Industries, Inc.

Subject Component \
0     Driveshaft Can Break POWER TRAIN
1     Charger Adapter May Cause Arcing or Shock Risk ELECTRICAL SYSTEM
2     Cooktop Burner Tube May Crack and Cause Gas Leak EQUIPMENT
3     Loss of Headlights and Taillights/FMVSS 108 ELECTRICAL SYSTEM
4     Spare Tire Carrier May Detach EQUIPMENT

Potentially Affected Recall Description \
0      18.0 GKN Automotive (GKN) is recalling certain repl...
1      130.0 N&B Mobility Solutions LLC (Nivion) is recalli...
2      396.0 Forest River, Inc. (Forest River) is recalling...
3      74469.0 Kia America, Inc. (Kia) is recalling certain 2...
4      107.0 Winnebago Industries, Inc. (Winnebago) is reca...

Consequence Summary \
0 A cracked or broken driveshaft can cause a los...
1 Inadequate clearance between DC busbars may ca...
2 A gas leak in the presence of an ignition sour...
3 A loss of headlights and taillights can reduce...
4 A detached spare tire carrier can become a roa...

Corrective Action Recall Type_Equipment \
0 GKN will reimburse the cost of a replacement d... True
1 Nivion will replace the defective adapters, fr... True
2 Owners are advised not to use the cooktop until... False
3 Dealers will update the BDC software, free of ... False
4 Dealers will inspect, replace, and correctly t... False

Recall Type_Tire Recall Type_Vehicle
0      False      False
1      False      False
2      False      True

```

```

▶ df.info()

→ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 28671 entries, 0 to 28670
Data columns (total 12 columns):
 #   Column           Non-Null Count Dtype
 ---  -- 
 0   Report Received Date    28671 non-null object
 1   NHTSA ID            28671 non-null object
 2   Manufacturer        28671 non-null object
 3   Subject             28671 non-null object
 4   Component           28671 non-null object
 5   Potentially Affected 28671 non-null float64
 6   Recall Description   28671 non-null object
 7   Consequence Summary 28671 non-null object
 8   Corrective Action    28671 non-null object
 9   Recall Type_Equipment 28671 non-null bool
 10  Recall Type_Tire     28671 non-null bool
 11  Recall Type_Vehicle  28671 non-null bool
dtypes: bool(3), float64(1), object(8)
memory usage: 2.1+ MB

```

6. Find out outliers (manually)

```

▶ import numpy as np

# Specify the column to analyze for outliers
col = 'Potentially Affected'

# Calculate Q1, Q3, and IQR
Q1 = df[col].quantile(0.25)
Q3 = df[col].quantile(0.75)
IQR = Q3 - Q1

# Define lower and upper bounds
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

# Identify outliers
outliers = df[(df[col] < lower_bound) | (df[col] > upper_bound)]

# Display the outliers
print(f"Outliers in '{col}':")
print(outliers)

```

Outliers in 'Potentially Affected':			
	Report Received Date	NHTSA ID	Manufacturer \
3	01/13/2025	25V006000	Kia America, Inc.
7	01/06/2025	25V002000	Tesla, Inc.
14	12/23/2024	24E110000	Horizon Global
21	12/20/2024	24V957000	Ford Motor Company
22	12/20/2024	24V954000	Ford Motor Company
...
28658	10/06/1966	66V004002	Ford Motor Company
28666	09/29/1966	66V003000	Honda (American Honda Motor Co.)
28668	01/19/1966	66V032001	General Motors, LLC
28669	01/19/1966	66V032003	General Motors, LLC
28670	01/19/1966	66V032004	General Motors, LLC
			Subject \
3			Loss of Headlights and Taillights/FMVSS 108
7			Rearview Camera Image May Fail/FMVSS 111
14			Tow Vehicle May Separate From Hitch Receiver Lock
21			High Pressure Fuel Pump May Fail
22			High Voltage Battery May Short Circuit

	Subject \	Component	Potentially Affected \
3	Loss of Headlights and Taillights/FMVSS 108	ELECTRICAL SYSTEM	74469.0
7	Rearview Camera Image May Fail/FMVSS 111	BACK OVER PREVENTION	239382.0
14	Tow Vehicle May Separate From Hitch Receiver Lock	TRAILER HITCHES	145431.0
21	High Pressure Fuel Pump May Fail	FUEL SYSTEM, DIESEL	295449.0
22	High Voltage Battery May Short Circuit	ELECTRICAL SYSTEM	20484.0
...
28658	Loss of Headlights and Taillights/FMVSS 108	SEAT BELTS	65000.0
28666	Rearview Camera Image May Fail/FMVSS 111	POWER TRAIN	18572.0
28668	Tow Vehicle May Separate From Hitch Receiver Lock	STEERING	138878.0
28669	High Pressure Fuel Pump May Fail	STEERING	70644.0
28670	High Voltage Battery May Short Circuit	STEERING	68184.0

Recall Description \			
3	Kia America, Inc. (Kia) is recalling certain 2...		
7	Tesla, Inc. (Tesla) is recalling certain 2024-...		
14	Horizon Global (Horizon) is recalling certain ...		
21	Ford Motor Company (Ford) is recalling certain...		
22	Ford Motor Company (Ford) is recalling certain...		
...	...		
28658		Not Known	
28666		Not Known	
28668		Not Known	
28669		Not Known	
28670		Not Known	
		Consequence Summary \	
3	A loss of headlights and taillights can reduce...		
7	A rearview camera that does not display an ima...		
14	A separated cap can allow the hitch to separat...		
21	High pressure Fuel pump failure can cause a lo...		
22	Battery failure can cause a loss of drive pow...		
...	...		
28658		Unknown	
28666		Unknown	
28668		Unknown	
28669		Unknown	
28670		Unknown	

	Corrective Action \
3	Dealers will update the BDC software, free of ...
7	Tesla released an over-the-air (OTA) software ...
14	Dealers will replace the hitch receiver locks,...
21	Dealers will update the powertrain control mod...
22	Dealers will perform a battery energy control ...
...	...
28658	Unknown
28666	Unknown
28668	Unknown
28669	Unknown
28670	Unknown

	Recall Type_Equipment	Recall Type_Tire	Recall Type_Vehicle
3	False	False	True
7	False	False	True
14	True	False	False
21	False	False	True
22	False	False	True
...
28658	False	False	True
28666	False	False	True
28668	False	False	True
28669	False	False	True
28670	False	False	True

[5063 rows x 12 columns]

7. standardization and normalization of column

```
▶ from sklearn.preprocessing import StandardScaler, MinMaxScaler
# Standardization: Transform data to have a mean of 0 and a standard deviation of 1
standard_scaler = StandardScaler()
df['Potentially Affected (Standardized)'] = standard_scaler.fit_transform(df[['Potentially Affected']])

# Normalization: Scale data between 0 and 1
min_max_scaler = MinMaxScaler()
df['Potentially Affected (Normalized)'] = min_max_scaler.fit_transform(df[['Potentially Affected']])

# Display the updated DataFrame
print(df[['Potentially Affected', 'Potentially Affected (Standardized)', 'Potentially Affected (Normalized)']].head())
```

	Potentially Affected	Potentially Affected (Standardized)	Potentially Affected (Normalized)
0	18.0	-0.122429	5.625000e-07
1	130.0	-0.122129	4.062500e-06
2	396.0	-0.121415	1.237500e-05
3	74469.0	0.077295	2.327156e-03
4	107.0	-0.122190	3.343750e-06

Conclusion:

This experiment demonstrated effective data cleaning and preparation techniques. Issues such as missing values, irrelevant data, and outliers were addressed, and the dataset was scaled for uniformity. These steps are essential for ensuring high-quality data and reliable model outcomes.

Experiment No:2

Aim: Data Visualization/ Exploratory data Analysis using Matplotlib and Seaborn.

Theory:

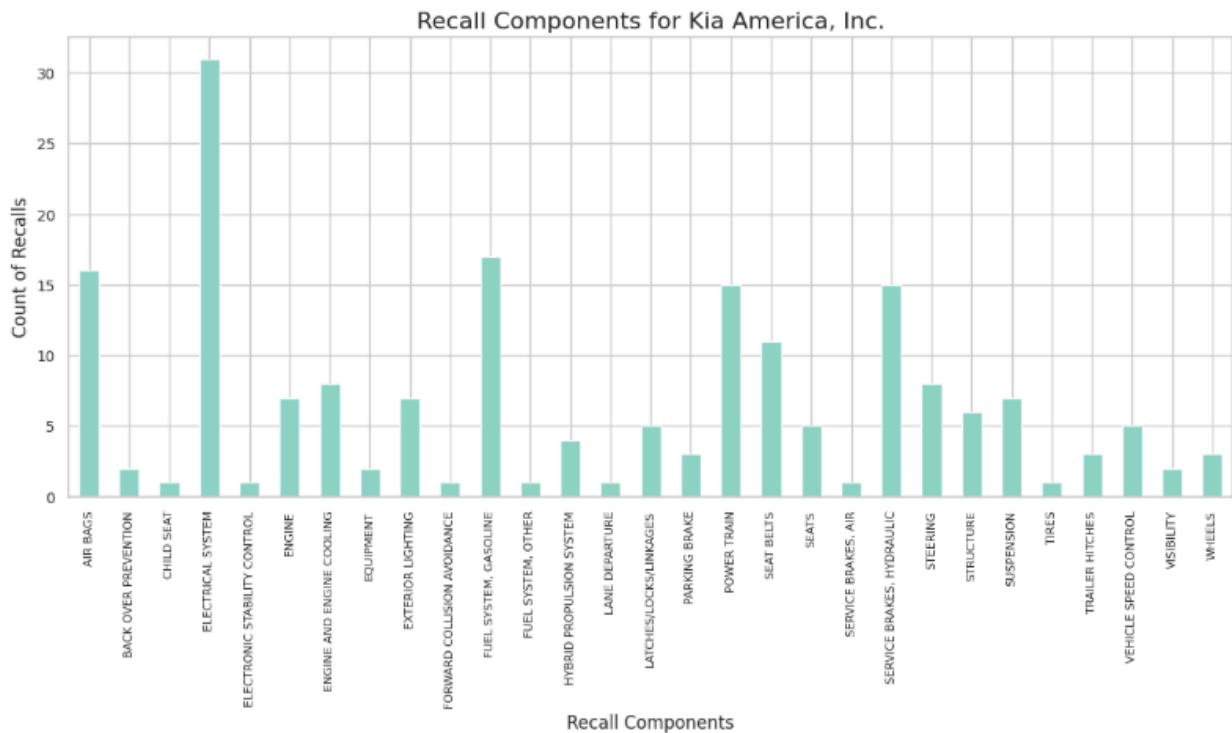
1. Bar graph & contingency table using any two features :

1. Bar Graph:

This graph visualizes the count of recalls for different components associated with Kia America, Inc. Each bar represents a specific component and the number of recalls it has. It helps identify which components have the most recalls, aiding in targeting problem areas.

```
import pandas as pd
import matplotlib.pyplot as plt

kia_df = df[df['Manufacturer'] == 'Kia America, Inc.']
crosstab_kia = pd.crosstab(kia_df['Component'], kia_df['Manufacturer'])
plt.figure(figsize=(12, 6))
crosstab_kia.plot.bar(figsize=(12, 6), colormap='Set3', rot=90)
plt.legend().set_visible(False)
plt.title('Recall Components for Kia America, Inc.', fontsize=16)
plt.xlabel('Recall Components', fontsize=12)
plt.ylabel('Count of Recalls', fontsize=12)
plt.xticks(fontsize=8, rotation=90)
plt.yticks(fontsize=10)
plt.tight_layout()
plt.subplots_adjust(bottom=0.2)
plt.show()
```



2. Contingency table: Displays the frequency of recalls for components against manufacturers. For Kia America, Inc., it shows the relationship between components and recall counts. Thus it helps in understanding how components vary by recall count across manufacturers.

```
import pandas as pd
import matplotlib.pyplot as plt

# Assuming 'df' is your DataFrame

# Create a contingency table (crosstab) for Component vs Manufacturer
crosstab = pd.crosstab(df['Component'], df['Manufacturer'])

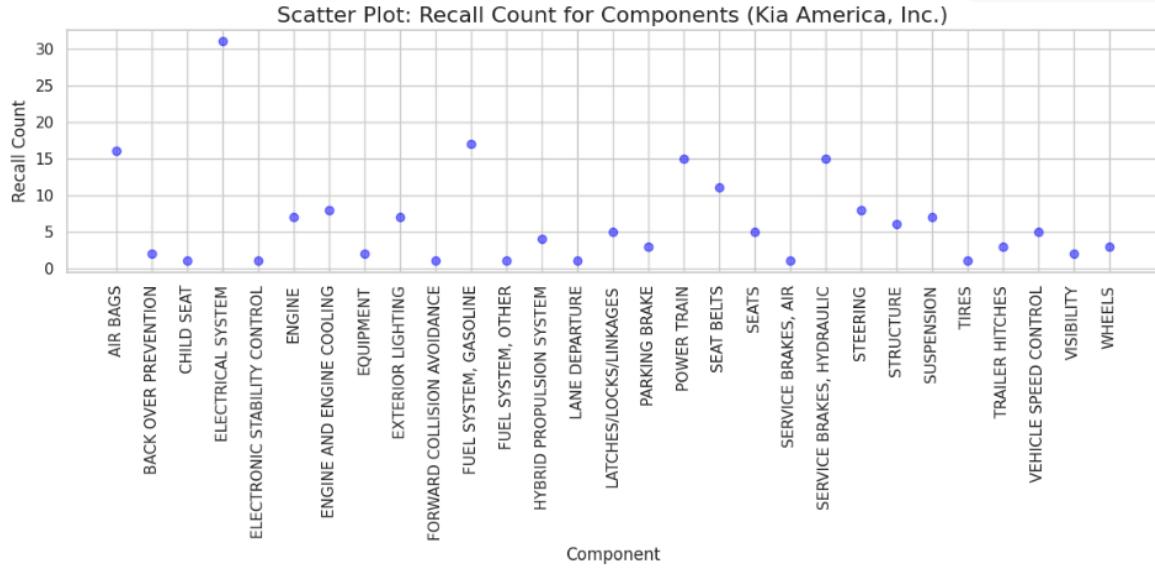
# Display the contingency table
print("Contingency Table:")
print(crosstab)

Contingency Table:
Manufacturer          1888653 Ontario Inc \
Component
AIR BAGS                      0
BACK OVER PREVENTION           0
CHILD SEAT                     0
COMMUNICATION                  0
ELECTRICAL SYSTEM               0
ELECTRONIC STABILITY CONTROL   0
ELECTRONIC STABILITY CONTROL (ESC) 0
ENGINE                         0
ENGINE AND ENGINE COOLING      0
EQUIPMENT                      1
EQUIPMENT ADAPTIVE/MOBILITY    0
EXTERIOR LIGHTING              0
FORWARD COLLISION AVOIDANCE    0
FUEL SYSTEM, DIESEL             0
FUEL SYSTEM, GASOLINE            0
FUEL SYSTEM, OTHER               0
HYBRID PROPULSION SYSTEM        0
INTERIOR LIGHTING               0
LANE DEPARTURE                  0
```

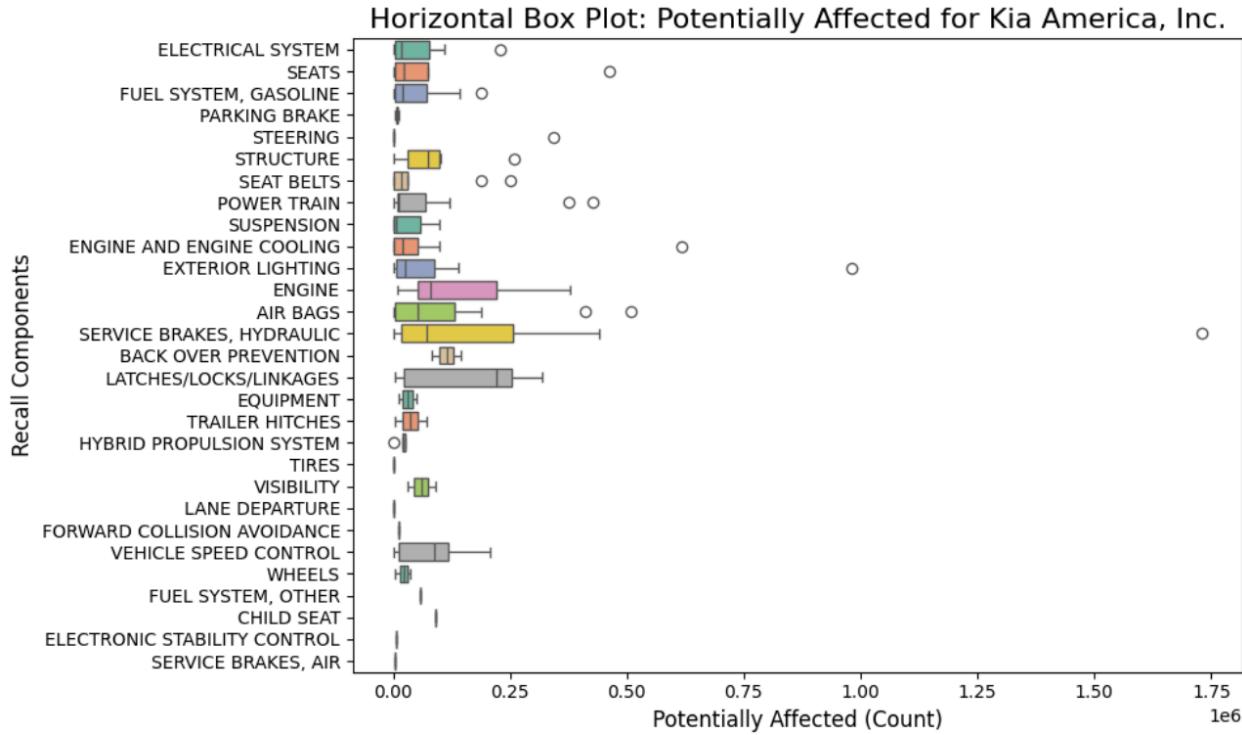
2. Plot Scatter plot, box plot, Heatmap using seaborn.

1. Scatter plot :This plot shows the recall count for various components of Kia America, Inc. Each point represents a component and its associated recall count. Thus it highlights the distribution and patterns of recall counts, identifying outliers or components with extreme values.

```
import pandas as pd
import matplotlib.pyplot as plt
kia_df = df[df['Manufacturer'] == 'Kia America, Inc.']
crosstab_kia = pd.crosstab(kia_df['Component'], kia_df['Manufacturer'])
scatter_data = crosstab_kia.reset_index()
scatter_data = scatter_data.melt(id_vars=['Component'], value_vars=crosstab_kia.columns, var_name='Manufacturer', value_name='Recall Count')
plt.figure(figsize=(12, 6))
plt.scatter(scatter_data['Component'], scatter_data['Recall Count'], c='blue', alpha=0.5)
plt.title('Scatter Plot: Recall Count for Components (Kia America, Inc.)', fontsize=16)
plt.xlabel('Component', fontsize=12)
plt.ylabel('Recall Count', fontsize=12)
plt.xticks(rotation=90)
plt.tight_layout()
plt.show()
```



2. Box plot: The box plot shows the distribution of the number of potentially affected vehicles for each recalled component. The horizontal layout makes it easier to compare components. It identifies the spread, central tendency, and outliers in the data, with whiskers representing the data range and dots as potential outliers.



```
#box plot
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

kia_df = df[df['Manufacturer'] == 'Kia America, Inc.']
plt.figure(figsize=(10, 6))
sns.boxplot(x='Potentially Affected', y='Component', data=kia_df, palette='Set2')
plt.title('Horizontal Box Plot: Potentially Affected for Kia America, Inc.', fontsize=16)
plt.xlabel('Potentially Affected (Count)', fontsize=12)
plt.ylabel('Recall Components', fontsize=12)
plt.xticks(fontsize=10)
plt.yticks(fontsize=10)

plt.tight_layout()
plt.show()
```

3. Heat map: The heatmap visualizes the contingency table data, where the intensity of color represents the recall count for each component. Thus helps in identifying which components have a higher recall count at a glance, making it easy to spot patterns or correlations.

```
# Load data (ensure this is defined)
df = pd.read_csv('Recalls_Data.csv') # Update with the correct file path

# Filter data for selected manufacturers
selected_manufacturers = []
    'Kia America, Inc.', 'Ford Motor Company', "Nissan North America, Inc.", 'Mercedes-Benz
[]

filtered_df = df[df['Manufacturer'].isin(selected_manufacturers)]

# Create crosstab for heatmap
heatmap_data = pd.crosstab(filtered_df['Component'], filtered_df['Manufacturer'])

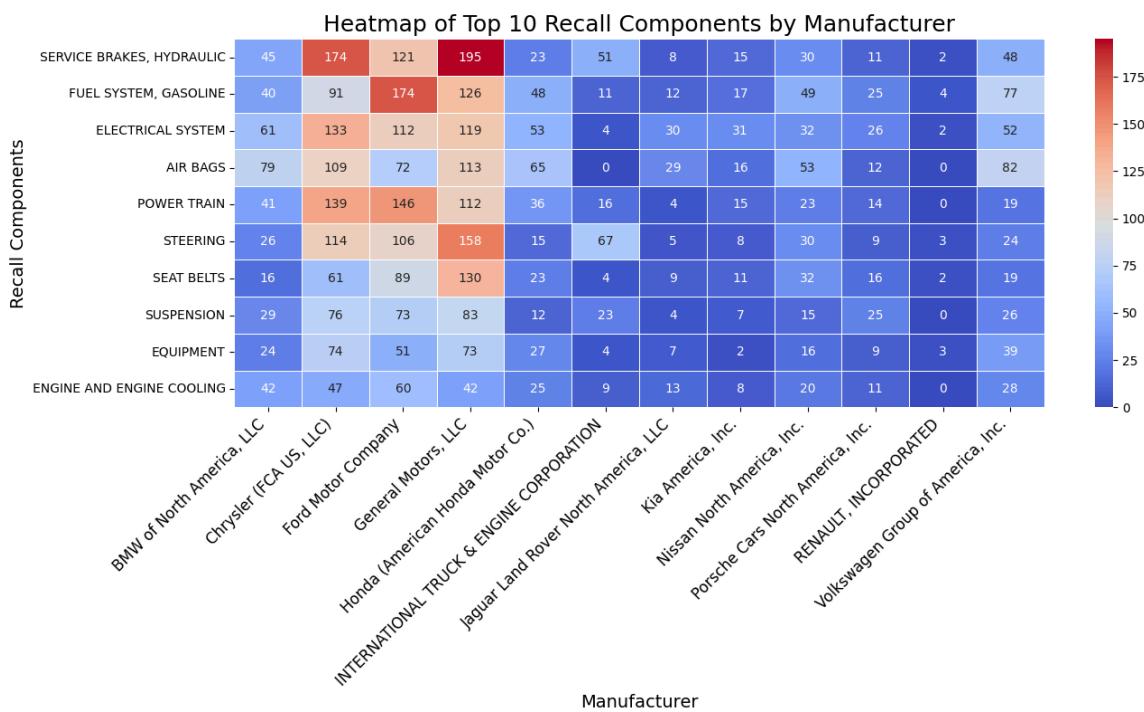
# Fill NaN values with 0
heatmap_data = heatmap_data.fillna(0)

# Sort by total recalls and select only the top 10 components
top_10_components = heatmap_data.sum(axis=1).nlargest(10).index
heatmap_data = heatmap_data.loc[top_10_components]

# Plot heatmap
plt.figure(figsize=(14, 8))
sns.heatmap(heatmap_data, annot=True, cmap='coolwarm', fmt='d', linewidths=0.5)

# Labels and title
plt.title('Heatmap of Top 10 Recall Components by Manufacturer', fontsize=18)
plt.xlabel('Manufacturer', fontsize=14)
plt.ylabel('Recall Components', fontsize=14)
plt.xticks(fontsize=12, rotation=45, ha="right")
plt.yticks(fontsize=10)

plt.tight_layout()
plt.show()
```

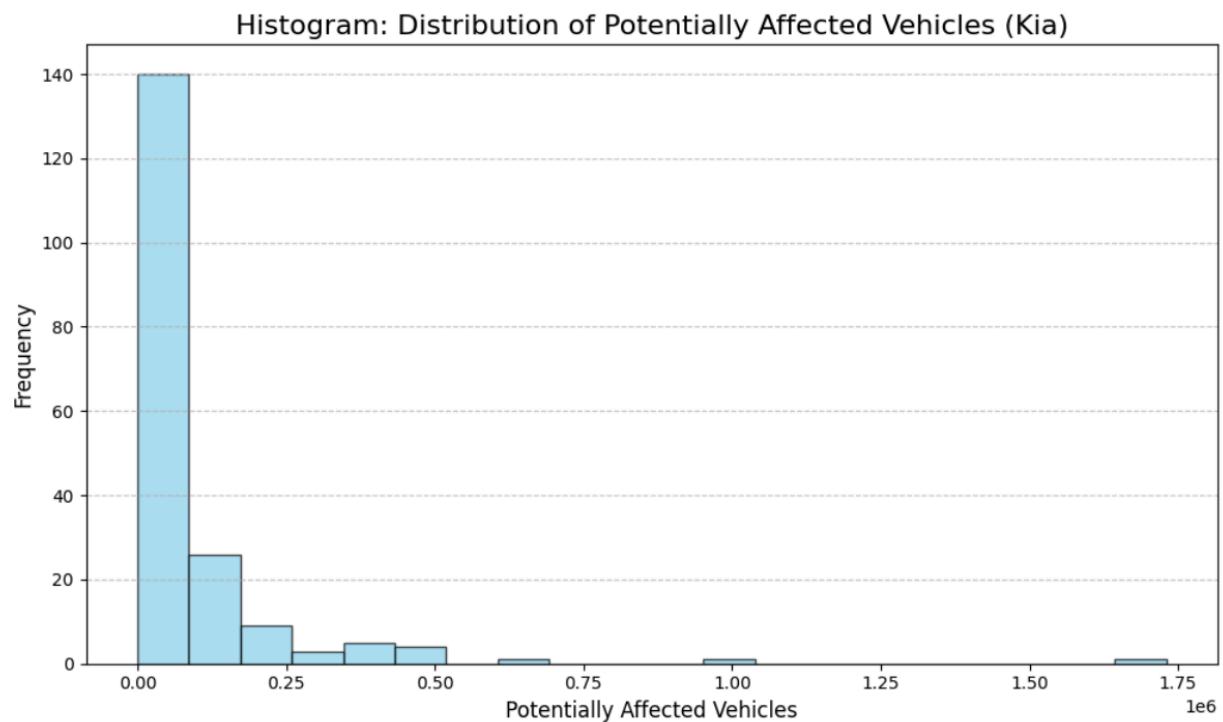


3. Create histogram and normalized Histogram.

The histogram shows the frequency distribution of the "Potentially Affected" vehicles.

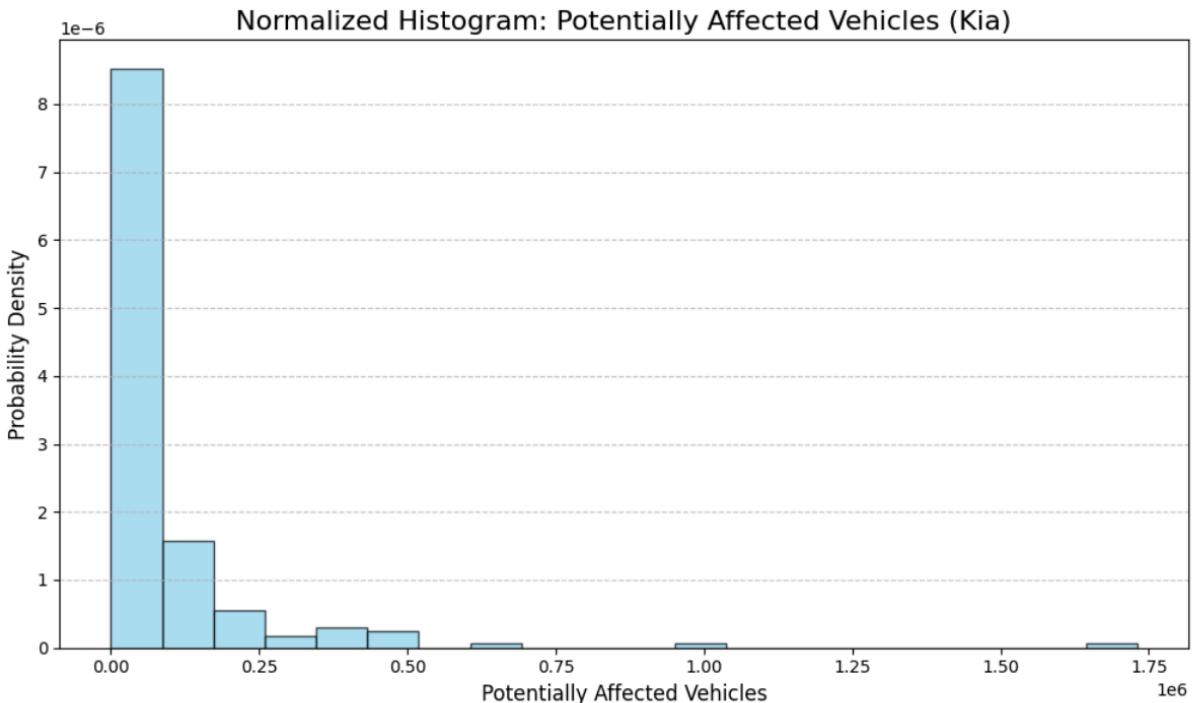
Regular Histogram: It displays counts in bins.

```
import pandas as pd, matplotlib.pyplot as plt
kia_df = df[df['Manufacturer'] == 'Kia America, Inc.']
plt.figure(figsize=(10,6))
plt.hist(kia_df['Potentially Affected'], bins=20, color='skyblue', edgecolor='black', alpha=0.7)
plt.title('Histogram: Distribution of Potentially Affected Vehicles (Kia)', fontsize=16)
plt.xlabel('Potentially Affected Vehicles', fontsize=12)
plt.ylabel('Frequency', fontsize=12)
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()
```



Normalized Histogram: It displays probability density.

```
import pandas as pd, matplotlib.pyplot as plt
kia_df = df[df['Manufacturer'] == 'Kia America, Inc.']
plt.figure(figsize=(10,6))
plt.hist(kia_df['Potentially Affected'], bins=20, color='skyblue', edgecolor='black', alpha=0.7, density=True)
plt.title('Normalized Histogram: Potentially Affected Vehicles (Kia)', fontsize=16)
plt.xlabel('Potentially Affected Vehicles', fontsize=12)
plt.ylabel('Probability Density', fontsize=12)
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()
```



This highlights the data's skewness, central tendency, and spread. The normalized version shows the proportion of each bin relative to the total.

4. Handle outlier using box plot and Inter quartile range.

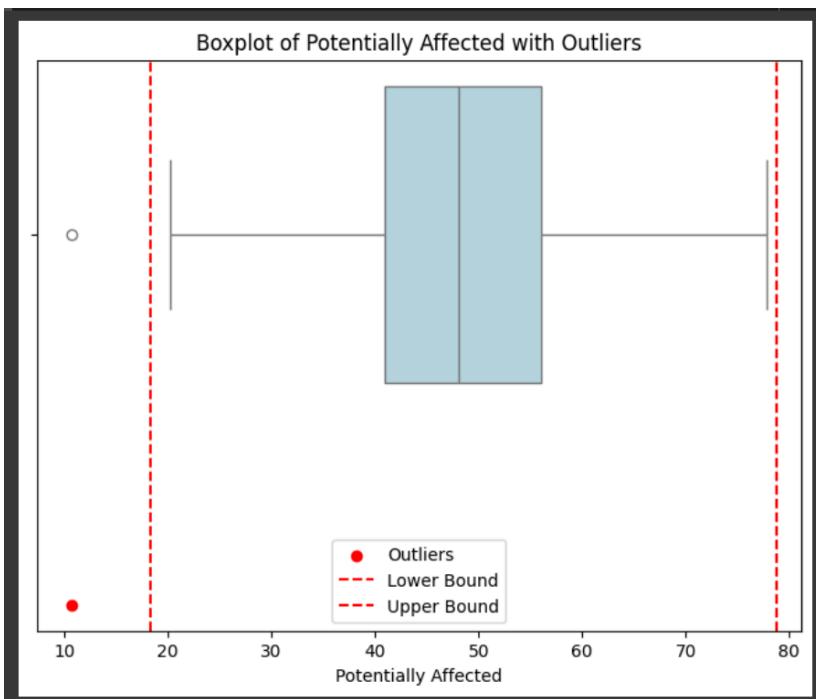
Outliers are values that deviate significantly from the rest of the data. Using the IQR method, data points beyond $Q1 - 1.5 \times IQR$ or $Q3 + 1.5 \times IQR$ are considered outliers.

Process: Outliers can be removed or replaced to reduce skewness and improve data accuracy.

```

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
np.random.seed(42)
df = pd.DataFrame({'Potentially Affected': np.random.normal(50, 15, 100)})
col = 'Potentially Affected'
Q1 = df[col].quantile(0.25)
Q3 = df[col].quantile(0.75)
IQR = Q3 - Q1
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR
outliers = df[(df[col] < lower_bound) | (df[col] > upper_bound)]
plt.figure(figsize=(8, 6))
sns.boxplot(x=df[col], color='lightblue')
plt.scatter(outliers[col], [1] * len(outliers), color='red', label='Outliers', zorder=2)
plt.axline(lower_bound, color='red', linestyle='dashed', label='Lower Bound')
plt.axline(upper_bound, color='red', linestyle='dashed', label='Upper Bound')
plt.title(f'Boxplot of {col} with Outliers')
plt.legend()
plt.show()

```



Conclusion: Bar Graph and Scatter Plot are suitable for identifying trends and comparisons. Heatmap is excellent for understanding correlations between variables. Box Plot provides insights into distributions and outliers. Normalized Histogram helps understand probabilities and densities. Outlier Handling improves the accuracy of the results.

Experiment 3

Aim: Perform Data Modeling on dataset.

Theory:

- a. Partition the data set, for example 75% of the records are included in the training data set and 25% are included in the test data set.

In this experiment, we partitioned a dataset of vehicle recall data into a training set (75%) and a test set (25%) and validated this partitioning using a two-sample Z-test. The dataset consists of 28,671 records with features such as manufacturer, recall components, and the number of potentially affected vehicles.

```
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
import scipy.stats as stats

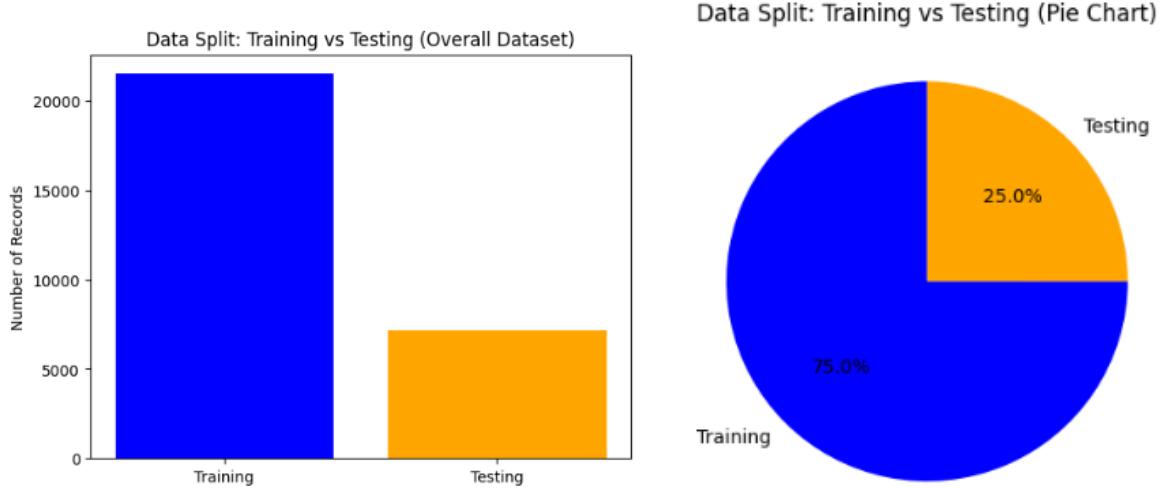
df=pd.read_csv('Recalls_Data.csv')
train_data, test_data = train_test_split(df, test_size=0.25, random_state=42)
labels = ['Training', 'Testing']
sizes = [len(train_data), len(test_data)]
plt.bar(labels, sizes, color=['blue', 'orange'])
plt.title("Data Split: Training vs Testing (Overall Dataset)")
plt.ylabel("Number of Records")
plt.show()
plt.pie(sizes, labels=labels, autopct='%1.1f%%', colors=['blue', 'orange'], startangle=90)
plt.title("Data Split: Training vs Testing (Pie Chart)")
plt.show()
print("Total records in the training data set:", len(train_data))
print("Total records in the testing data set:", len(test_data))
```

- b. Use a bar graph and other relevant graph to confirm your proportions.

To confirm the proportions of the data split into training and test sets, we have used a bar graph and a pie chart:

Bar Graph: It displays the number of records in the training and test sets, allowing us to visually confirm the 75%-25% split. The x-axis represents the two sets, and the y-axis shows their respective record counts.

Pie Chart: This chart visually illustrates the percentage split between the training (75%) and test (25%) sets, providing an easy confirmation of the partition.



c. Identify the total number of records in the training data set.

We used the `train_test_split` method to split the dataset into training and test sets. The partition was visually confirmed through a bar plot, showing the expected 75%-25% split, with 21,503 records in the training set and 7,168 in the test set.\

Total records in the training data set: 21503

Total records in the testing data set: 7168

d. Validate partition by performing a two-sample Z-test.

To validate the partitioning, we performed a two-sample Z-test manually to compare the "Potentially Affected" values between the training and test datasets. The Z-statistic was calculated based on the means, standard deviations, and sizes of both samples. We then compared the calculated p-value to the significance level of 0.05.

The Z-test showed that there was no significant difference between the training and test datasets, as the p-value was greater than 0.05. This confirmed that the partitioning process was unbiased.

Conclusion:

The partitioning of the dataset into training and test sets was validated successfully. The Z-test confirmed that there was no significant difference between the datasets, making the partition reliable for further analysis.

Experiment 4

Aim: Implementation of Statistical Hypothesis Test using Scipy and Sci-kit learn.

Theory:

1. **Pearson's Correlation:** Pearson's correlation measures the linear relationship between two continuous variables. A coefficient close to 1 or -1 indicates a strong linear relationship, while a value near 0 suggests no linear association.

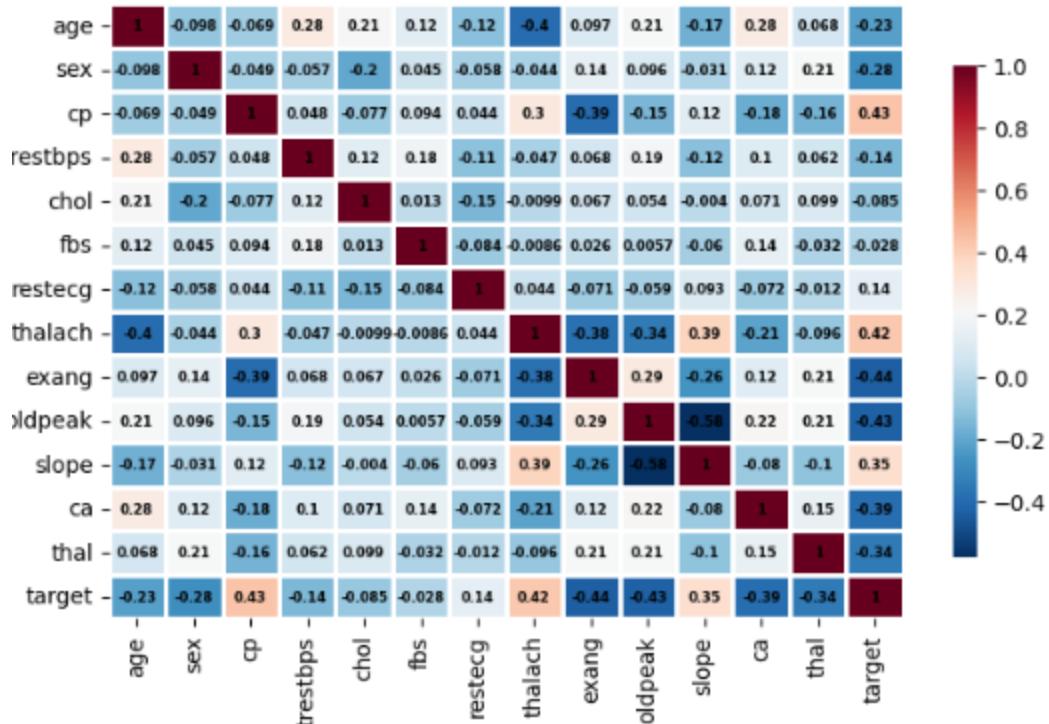
	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
age	1.000000	-0.098447	-0.068653	0.279351	0.213678	0.121308	-0.116211	-0.398522	0.096801	0.210013	-0.168814	0.276326	0.068001	-0.225439
sex	-0.098447	1.000000	-0.049353	-0.056769	-0.197912	0.045032	-0.058196	-0.044020	0.141664	0.096093	-0.030711	0.118261	0.210041	-0.280937
cp	-0.068653	-0.049353	1.000000	0.047608	-0.076904	0.094444	0.044421	0.295762	-0.394280	-0.149230	0.119717	-0.181053	-0.161736	0.433798
trestbps	0.279351	-0.056769	0.047608	1.000000	0.123174	0.177531	-0.114103	-0.046698	0.067616	0.193216	-0.121475	0.101389	0.062210	-0.144931
chol	0.213678	-0.197912	-0.076904	0.123174	1.000000	0.013294	-0.151040	-0.009940	0.067023	0.053952	-0.004038	0.070511	0.098803	-0.085239
fbs	0.121308	0.045032	0.094444	0.177531	0.013294	1.000000	-0.084189	-0.008567	0.025665	0.005747	-0.059894	0.137979	-0.032019	-0.028046
restecg	-0.116211	-0.058196	0.044421	-0.114103	-0.151040	-0.084189	1.000000	0.044123	-0.070733	-0.058770	0.093045	-0.072042	-0.011981	0.137230
thalach	-0.398522	-0.044020	0.295762	-0.046698	-0.009940	-0.008567	0.044123	1.000000	-0.378812	-0.344187	0.386784	-0.213177	-0.096439	0.421741
exang	0.096801	0.141664	-0.394280	0.067616	0.067023	0.025665	-0.070733	-0.378812	1.000000	0.288223	-0.257748	0.115739	0.206754	-0.436757
oldpeak	0.210013	0.096093	-0.149230	0.193216	0.053952	0.005747	-0.058770	-0.344187	0.288223	1.000000	-0.577537	0.222682	0.210244	-0.430696
slope	-0.168814	-0.030711	0.119717	-0.121475	-0.004038	-0.059894	0.093045	0.386784	-0.257748	-0.577537	1.000000	-0.080155	-0.104764	0.345877
ca	0.276326	0.118261	-0.181053	0.101389	0.070511	0.137979	-0.072042	-0.213177	0.115739	0.222682	-0.080155	1.000000	0.151832	-0.391724
thal	0.068001	0.210041	-0.161736	0.062210	0.098803	-0.032019	-0.011981	-0.096439	0.206754	0.210244	-0.104764	0.151832	1.000000	-0.344029
target	-0.225439	-0.280937	0.433798	-0.144931	-0.085239	-0.028046	0.137230	0.421741	-0.436757	-0.430696	0.345877	-0.391724	-0.344029	1.000000

```
import seaborn as sb
import matplotlib.pyplot as plt

plt.figure(figsize=(8, 5))

sb.heatmap(pearsoncorr,
           xticklabels=pearsoncorr.columns,
           yticklabels=pearsoncorr.columns,
           cmap='RdBu_r',
           annot=True,
           annot_kws={"size": 6, "weight": "bold", "color": "black"},
           linewidth=2,
           cbar_kws={"shrink": 0.8})

plt.show()
```



Result: There is a moderate positive relationship between cp and heart disease (target), with a correlation of 0.43.

2. **Spearman's Rank Correlation:** Spearman's rank correlation assesses the monotonic relationship between variables, relying on their ranks rather than raw data. It is suitable for ordinal or non-linear relationships.

```

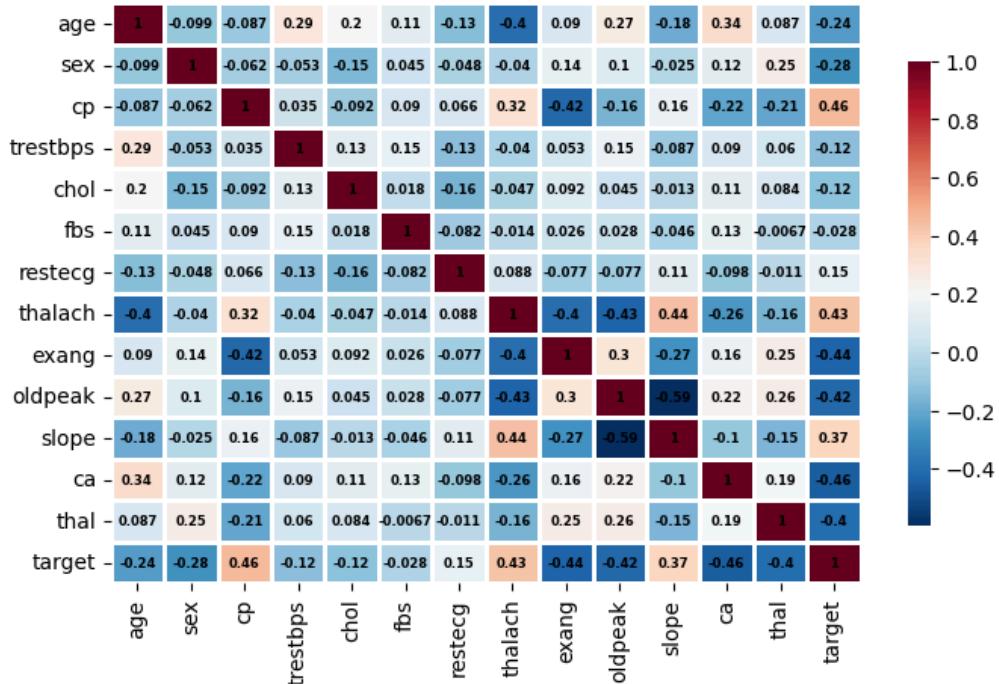
spearmancorr = df.corr(method='spearman')

plt.figure(figsize=(8, 5))

sb.heatmap(spearmancorr,
            xticklabels=spearmancorr.columns,
            yticklabels=spearmancorr.columns,
            cmap='RdBu_r',
            annot=True,
            annot_kws={"size": 6, "weight": "bold", "color": "black"},
            linewidth=2,
            cbar_kws={"shrink": 0.8})

plt.show()

```



Results: The correlation between cp (chest pain type) and target is 0.46, indicating a moderate positive association.

3. **Kendall's Rank Correlation:** Kendall's Tau measures the strength of the ordinal relationship between two variables by comparing the ranks of pairs. It is more robust to ties in data.

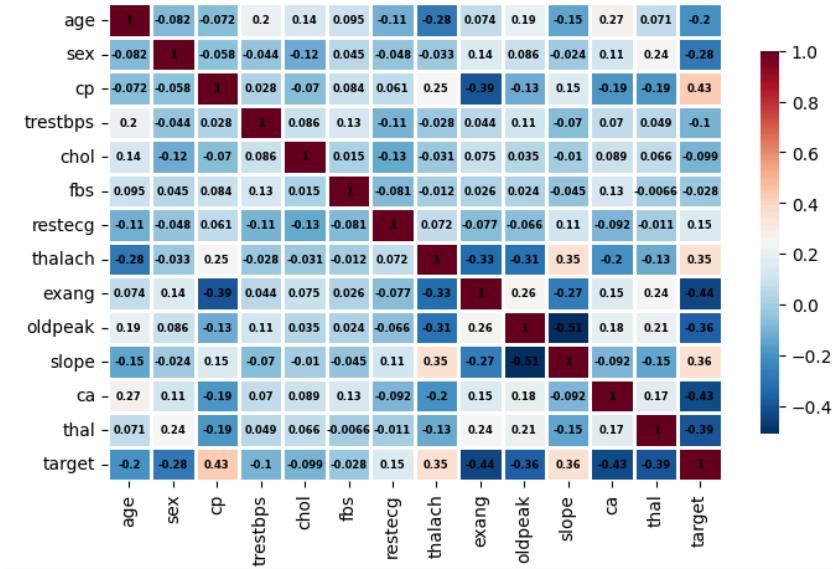
```
from scipy.stats import pearsonr, spearmanr, kendalltau, chi2_contingency

kendallcorr = df.corr(method='kendall')

plt.figure(figsize=(8, 5))

sb.heatmap(kendallcorr,
            xticklabels=kendallcorr.columns,
            yticklabels=kendallcorr.columns,
            cmap='RdBu_r',
            annot=True,
            annot_kws={"size": 6, "weight": "bold", "color": "black"},
            linewidth=2,
            cbar_kws={"shrink": 0.8})

plt.show()
```



Results: cp and target have a Kendall's Tau of 0.43, showing a positive relationship. target and thalach show a weaker but still positive correlation (0.35), indicating heart rate's relevance in predicting heart disease.

4. **Chi-Squared Test:** The Chi-Squared test assesses the independence of two categorical variables. A significant p-value indicates that the variables are dependent (associated).

```
import pandas as pd
from scipy.stats import chi2_contingency

contingency_table = pd.crosstab(df['sex'], df['target'])

chi2, p, dof, expected = chi2_contingency(contingency_table)

print("Chi-Squared Statistic:", chi2)
print("P-value:", p)
print("Degrees of Freedom:", dof)
print("Expected frequencies table:")
print(expected)

if p < 0.05:
    print("There is a significant association between the variables (reject the null hypothesis).")
else:
    print("There is no significant association between the variables (fail to reject the null hypothesis).")
```

```
Chi-Squared Statistic: 22.717227046576355
P-value: 1.8767776216941503e-06
Degrees of Freedom: 1
Expected frequencies table:
[[ 43.72277228  52.27722772]
 [ 94.27722772 112.72277228]]
There is a significant association between the variables (reject the null hypothesis).
```

Result: The Chi-Squared statistic of 22.72 (p-value = 1.88e-06) indicates a significant association between categorical variables (e.g., sex and target), suggesting their role in heart disease prediction.

Conclusion:

In this analysis, four statistical tests were applied to assess the relationships between various features and heart disease:

1. Pearson's correlation showed a moderate positive relationship between cp and heart disease (target), with a correlation of 0.43.
2. Spearman's rank correlation confirmed a moderate positive relationship between cp (chest pain type) and target (0.46).
3. Kendall's Tau also revealed a moderate positive association between cp and target (0.43)
4. The Chi-Squared test showed a significant association between categorical variables, with a Chi-Squared statistic of 22.72 and a p-value of 1.88e-06. Since the p-value of 1.88e-06 is much smaller than the commonly used significance level of 0.05, we reject the null hypothesis.

Experiment 5

Aim: Perform Regression Analysis using Scipy and Sci-kit learn.

Theory:

1. Linear Regression:

Linear Regression is a supervised learning algorithm used for predicting continuous numerical values. It assumes a linear relationship between the independent variables (features) and the dependent variable (target).

The model minimizes the sum of squared residuals (differences between actual and predicted values) to find the best-fitting line.

It is sensitive to outliers and works best when the data follows a linear trend. Performance is evaluated using metrics like Mean Squared Error (MSE) and R-squared (R^2).

2. Logistic Regression:

Logistic Regression is a classification algorithm used for predicting binary outcomes (0 or 1).

It applies the logistic (sigmoid) function to transform linear outputs into probabilities.

The model predicts a class based on a threshold, typically 0.5.

It uses optimization techniques like gradient descent to minimize the loss function (log loss).

Performance is assessed using accuracy, precision, recall, F1-score, and confusion matrices.

Implementation:

1. Import necessary libraries and load the dataset

```
[ ] import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, LogisticRegression
from sklearn.preprocessing import StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.metrics import mean_squared_error, r2_score, accuracy_score
```

```
[ ] df = pd.read_csv('diabetes_dataset_with_notes.csv')

[ ] df.head()

year gender age location race:AfricanAmerican race:Asian race:Caucasian race:Hispanic race:Other hypertension
0 2020 Female 32.0 Alabama 0 0 0 0 1 0
1 2015 Female 29.0 Alabama 0 1 0 0 0 0
2 2015 Male 18.0 Alabama 0 0 0 0 1 0
3 2015 Male 41.0 Alabama 0 0 1 0 0 0
4 2016 Female 52.0 Alabama 1 0 0 0 0 0

df.describe()
```

	year	age	race:AfricanAmerican	race:Asian	race:Caucasian
count	100000.000000	100000.000000	100000.000000	100000.000000	100000.000000
mean	2018.360820	41.885856	0.202230	0.200150	0.198760
std	1.345239	22.516840	0.401665	0.400114	0.399069
min	2015.000000	0.080000	0.000000	0.000000	0.000000
25%	2019.000000	24.000000	0.000000	0.000000	0.000000
50%	2019.000000	43.000000	0.000000	0.000000	0.000000
75%	2019.000000	60.000000	0.000000	0.000000	0.000000
max	2022.000000	80.000000	1.000000	1.000000	1.000000

2. Check for missing values and anomalies

```
print("\nMissing values after handling:\n", X.isnull().sum())

Missing values after handling:
age 0
gender 0
bmi 0
hbA1c_level 0
hypertension 0
heart_disease 0
bmi_age_interaction 0
hbA1c_blood_glucose 0
dtype: int64

numeric_columns = df.select_dtypes(include=[np.number]).columns
imputer = SimpleImputer(strategy="median")
df[numeric_columns] = imputer.fit_transform(df[numeric_columns])

df["bmi_age_interaction"] = df["bmi"] * df["age"]
df["hbA1c_blood_glucose"] = df["hbA1c_level"] * df["blood_glucose_level"]

df = df[df["blood_glucose_level"] < df["blood_glucose_level"].quantile(0.99)]
```

3. Choose your columns and convert categorical values into numerical values for better classification

```
df["gender"] = df["gender"].map({"Male": 0, "Female": 1})
df["smoking_history"] = df["smoking_history"].astype('category').cat.codes

features = ["age", "gender", "bmi", "hbA1c_level", "hypertension", "heart_disease",
            "bmi_age_interaction", "hbA1c_blood_glucose"]

X = df[features]
y_reg = df["blood_glucose_level"] # Target for Linear Regression
y_clf = df["diabetes"] # Target for Logistic Regression (Assuming it's binary 0/1)

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

4. Train dataset using linear regression

```
[35]
df = df.drop(columns=["location", "clinical_notes", "smoking_history", "gender"])

# Define features (X) and target (y)
X = df.drop(columns=["diabetes"])
y = df["diabetes"]

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize and train linear regression model
model = LinearRegression()
model.fit(X_train, y_train)

# Predictions
y_pred = model.predict(X_test)

# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"Mean Squared Error: {mse}")
print(f"R-squared Score: {r2}")
```

5. Finding out Mean Squared error and goodness of fit

```
Mean Squared Error: 0.041312657154867635
R-squared Score: 0.25314264247732976
```

6. Using Logistic Regression, find out accuracy score

```
X = df.drop(columns=["diabetes"])
y = df["diabetes"]

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize and train logistic regression model
model = LogisticRegression(max_iter=1000)
model.fit(X_train, y_train)

# Predictions
y_pred = model.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
report = classification_report(y_test, y_pred)

print(f"Accuracy: {accuracy}")
print("Classification Report:")
print(report)
```

Accuracy Score:

```
Accuracy: 0.962550087331758
Classification Report:
precision    recall   f1-score   support
      0.0       0.97      0.99      0.98     18322
      1.0       0.82      0.47      0.59      1144
      accuracy                           0.96     19466
      macro avg       0.89      0.73      0.79     19466
      weighted avg     0.96      0.96      0.96     19466
```

Conclusion:

Linear Regression Results:

- The Mean Squared Error (MSE) of 0.0413 indicates that the model's predictions are relatively close to the actual values.

- The R-squared score of 0.2531 shows that the independent variables explain about 25% of the variance in diabetes presence.

Logistic Regression Results:

- The model achieved an accuracy of 96.25%.
- The classification report shows that it performs well for non-diabetic cases (0) with high precision and recall.
- For diabetic cases (1), the recall is 0.47, meaning that while many diabetic cases are correctly identified, some are still being misclassified.

Experiment No: 6

Aim: Perform Classification modelling

a. Choose a classifier for classification problems.

b. Evaluate the performance of the classifier.

Perform Classification using the below 4 classifiers on the same dataset:

1. K-Nearest Neighbors (KNN)
2. Naive Bayes
3. Support Vector Machines (SVMs)
4. Decision Tree

Theory:

1. **Decision tree:** A Decision Tree is a supervised learning algorithm for classification and regression that splits data based on feature values to form a tree-like structure. It uses Gini Index or Entropy to determine splits and can be pruned to prevent overfitting. While easy to interpret and handle both numerical and categorical data, it can be unstable and prone to overfitting.

```
X = df[features]
y = df["diabetes"]
X_clean = X.copy()
X_clean.replace([np.inf, -np.inf], np.nan, inplace=True)
X_clean.dropna(axis=1, how='all', inplace=True)
X_clean.dropna(axis=0, how='all', inplace=True)
X_clean = X_clean.apply(lambda col: col.fillna(col.mean()), axis=0)
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X_clean)
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)
dt_model = DecisionTreeClassifier(random_state=42)
dt_model.fit(X_train, y_train)
y_pred_dt = dt_model.predict(X_test)
accuracy_dt = accuracy_score(y_test, y_pred_dt)
print("\nDecision Tree Classifier Results:")
print(f"Accuracy: {accuracy_dt:.2f}")
print("Classification Report for Decision Tree:")
print(classification_report(y_test, y_pred_dt))

cm = confusion_matrix(y_test, y_pred_dt)

print("Confusion Matrix:")
print(cm)
```

1. The Decision Tree classifier achieved **95% accuracy**, indicating strong overall performance. However, a closer look at the classification report and confusion matrix reveals an imbalance in predictive capability between the two classes:

- **Class 0 (Non-Diabetic):** The model performed exceptionally well, with **98% precision** and **97% recall**, meaning most non-diabetic cases were correctly classified.
- **Class 1 (Diabetic):** The model struggled with diabetic cases, achieving **67% precision** and **70% recall**, meaning **30% of actual diabetic cases were misclassified as non-diabetic**.

2. From the **confusion matrix**, we observe:

- **True Negatives (TN):** 17,801 (correctly classified non-diabetic cases)
- **False Positives (FP):** 495 (misclassified non-diabetic cases as diabetic)
- **False Negatives (FN):** 421 (misclassified diabetic cases as non-diabetic)
- **True Positives (TP):** 1,003 (correctly classified diabetic cases)

Decision Tree Classifier Results:

Accuracy: 0.95

Classification Report for Decision Tree:

	precision	recall	f1-score	support
0	0.98	0.97	0.97	18296
1	0.67	0.70	0.69	1424
accuracy			0.95	19720
macro avg	0.82	0.84	0.83	19720
weighted avg	0.95	0.95	0.95	19720

Confusion Matrix:

```
[[17801  495]
 [ 421 1003]]
```

3. While the model performs well overall, the lower recall for diabetic cases indicates that a significant number of diabetic patients are not being correctly identified. This could be due to **class imbalance** in the dataset, where non-diabetic cases dominate.

2. **Support Vector Machine (SVM):** A Support Vector Machine (SVM) finds the optimal hyperplane to separate classes, using support vectors to maximize the margin. It employs the kernel trick (Linear, Polynomial, RBF) for non-linearly

separable data. SVMs perform well on high-dimensional data and small datasets but can be computationally expensive and require careful kernel selection.

```
X_clean = X.copy()
X_clean.replace([np.inf, -np.inf], np.nan, inplace=True)
X_clean.dropna(axis=1, how='all', inplace=True)
X_clean.dropna(axis=0, how='all', inplace=True)
X_clean = X_clean.apply(lambda col: col.fillna(col.mean()), axis=0)
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X_clean)
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)
svm_model = SVC(kernel='linear', random_state=42)
svm_model.fit(X_train, y_train)
y_pred_svm = svm_model.predict(X_test)
accuracy_svm = accuracy_score(y_test, y_pred_svm)
print("\nSupport Vector Machine (SVM) Classifier Results:")
print(f"Accuracy: {accuracy_svm:.2f}")
print("Classification Report for SVM:")
print(classification_report(y_test, y_pred_svm))
```

1. The Support Vector Machine (SVM) classifier achieved **96% accuracy**, indicating strong overall performance. However, a closer look at the classification report and confusion matrix reveals an imbalance in predictive capability between the two classes:

- **Class 0 (Non-Diabetic):** The model performed exceptionally well, with **96% precision and 100% recall**, meaning almost all non-diabetic cases were correctly classified.
- **Class 1 (Diabetic):** The model struggled with diabetic cases, achieving **92% precision and 48% recall**, meaning **52% of actual diabetic cases were misclassified as non-diabetic**.

2. From the confusion matrix, we observe:

- **True Negatives (TN):** 17,801 (correctly classified non-diabetic cases)
- **False Positives (FP):** 495 (misclassified non-diabetic cases as diabetic)
- **False Negatives (FN):** 421 (misclassified diabetic cases as non-diabetic)
- **True Positives (TP):** 1,003 (correctly classified diabetic cases)

```

Support Vector Machine (SVM) Classifier Results:
Accuracy: 0.96
Classification Report for SVM:
precision    recall   f1-score   support
          0       0.96      1.00      0.98     18296
          1       0.92      0.48      0.64     1424
                                              accuracy      0.96     19720
                                              macro avg      0.94      0.74      0.81     19720
                                              weighted avg      0.96      0.96      0.95     19720

Confusion Matrix:
[[17801  495]
 [ 421 1003]]

```

3. While the model performs well overall, the **low recall for diabetic cases** indicates that a significant number of diabetic patients are not being correctly identified. This could be due to class imbalance in the dataset, where non-diabetic cases dominate.

Conclusion: In this experiment, we implemented two classification models—Support Vector Machines (SVM), and Decision Tree—on the diabetes dataset to evaluate their performance.

- Decision Tree achieved high accuracy (95%) but showed a bias towards non-diabetic cases due to class imbalance. It had excellent precision and recall for non-diabetic cases but struggled with diabetic cases.
- Support Vector Machine (SVM) performed slightly better overall (96% accuracy) but had low recall (48%) for diabetic cases, meaning it misclassified many diabetic patients.

EXPERIMENT NO:7

Aim: To implement different clustering algorithms.

Problem Statement: a) Clustering algorithm for unsupervised classification (K-means, density based (DBSCAN), Hierarchical clustering)
b) Plot the cluster data and show mathematical steps.

Theory:

1. Importing Libraries

```
▶ import pandas as pd
  import numpy as np
  from sklearn.cluster import KMeans
  from sklearn.preprocessing import StandardScaler
  from sklearn.decomposition import PCA
  from sklearn.preprocessing import LabelEncoder
  from sklearn.cluster import DBSCAN
  import scipy.cluster.hierarchy as sch
  import matplotlib.pyplot as plt
  import seaborn as sns
```

2. Importing dataset

```
df=pd.read_csv('diabetes_dataset_with_notes.csv')
df = pd.DataFrame(df)
```

3. Transformation

```
label_encoder = LabelEncoder()
df['gender_encoded'] = label_encoder.fit_transform(df['gender'])
df['smoking_history_encoded'] = label_encoder.fit_transform(df['smoking_history'])

features = ['age', 'race:AfricanAmerican', 'race:Asian', 'race:Caucasian', 'race:Hispanic',
            'race:Other', 'hypertension', 'heart_disease', 'smoking_history_encoded', 'bmi',
            'hbA1c_level', 'blood_glucose_level']

scaler = StandardScaler()
df_scaled = scaler.fit_transform(df[features])
```

4. **K Means** :-K-Means Clustering: A centroid-based clustering algorithm that partitions data into k clusters by minimizing the distance between data points and their respective cluster centers. It works well with well-separated clusters but assumes a spherical shape.

```
kmeans = KMeans(n_clusters=3, random_state=42)
df['kmeans_cluster'] = kmeans.fit_predict(df_scaled)

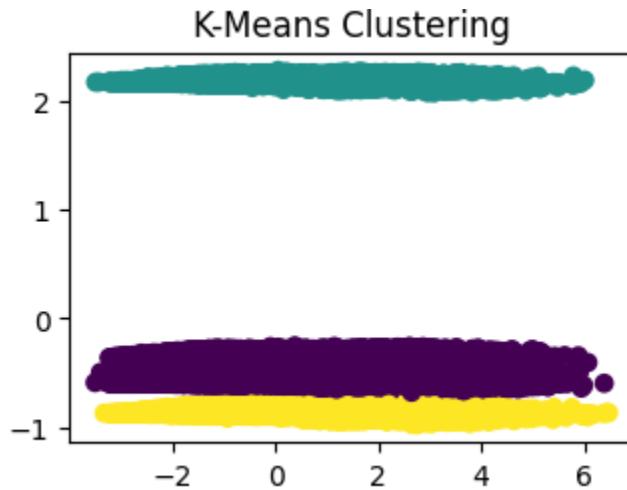
pca = PCA(n_components=2)
df_pca = pca.fit_transform(df_scaled)

plt.subplot(2, 2, 1)
plt.scatter(df_pca[:, 0], df_pca[:, 1], c=df['kmeans_cluster'], cmap='viridis')
plt.title("K-Means Clustering")

plt.tight_layout()
plt.show()
```

This code performs K-Means clustering on a dataset with 3 clusters, assigns cluster labels, and visualizes the results using PCA (Principal Component Analysis) for dimensionality reduction.

- **K-Means Clustering:** It fits the scaled dataset and assigns each data point a cluster label (0, 1, or 2), which is stored in `df['kmeans_cluster']`.
- **PCA Transformation:** The dataset is reduced to 2D for visualization.
- **Plotting:** A scatter plot is created using the first two PCA components, with points colored based on their assigned clusters.



- Successfully identified three distinct clusters.
- The clusters are horizontally aligned, suggesting that one or two dominant features are driving the segmentation.
- Works well for structured, well-separated data but may struggle with complex patterns.

5. **DBSCAN Clustering:-**DBSCAN (Density-Based Spatial Clustering of Applications with Noise): A clustering algorithm that groups data points based on density rather than predefined clusters. It can detect arbitrarily shaped clusters and outliers but is sensitive to parameter selection (eps, min_samples).

```
dbscan = DBSCAN(eps=0.5, min_samples=2)
df['dbscan_cluster'] = dbscan.fit_predict(df_scaled)

plt.subplot(2, 2, 2)
plt.scatter(df_pca[:, 0], df_pca[:, 1], c=df['dbscan_cluster'], cmap='plasma')
plt.title("DBSCAN Clustering")

plt.tight_layout()
plt.show()
```

This code applies DBSCAN clustering on a dataset and visualizes the results using PCA for dimensionality reduction.

Steps:

1. DBSCAN Clustering:

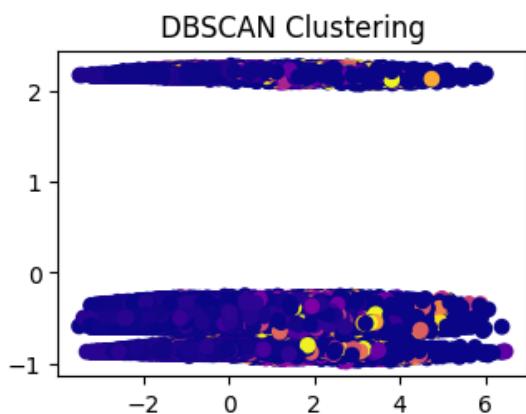
- DBSCAN(eps=0.5, min_samples=2): Groups points based on density.
- Noise points (outliers) are labeled as -1.
- Cluster labels are stored in df['dbscan_cluster'].

2. PCA for Visualization:

- Reduces data to 2D for plotting.

3. Scatter Plot:

- Points are colored by their DBSCAN cluster labels using the 'plasma' colormap.
- Outliers appear as separate scattered points.



- Identified dense regions, but the separation is not as clear as K-Means.
- Some points were classified as noise (outliers).
- Suitable for non-linear and arbitrarily shaped clusters, but parameter tuning (eps, min_samples) is crucial for better results.

6. Hierarchical clustering

```
# Take a random sample of 5000 to reduce memory usage
sample_size = 5000
df_sample = df.sample(n=sample_size, random_state=42)
data_sample = df_sample[numeric_features]

# Standardize sample data
scaler = StandardScaler()
data_sample_scaled = scaler.fit_transform(data_sample)

# Apply Hierarchical Clustering on the sample
hc = AgglomerativeClustering(n_clusters=3, linkage='ward')
labels = hc.fit_predict(data_sample_scaled) # Store labels separately

df_sample['hc_cluster'] = labels # Assign cluster labels to the sample

# Scatter plot of clusters
plt.scatter(data_sample_scaled[:, 0], data_sample_scaled[:, 1], c=df_sample['hc_cluster'], cmap='plasma')
plt.title("Hierarchical Clustering (Sampled Data)")
plt.xlabel(numeric_features[0])
plt.ylabel(numeric_features[1])
plt.colorbar(label='Cluster')
plt.show()
```

1. Data Preprocessing:

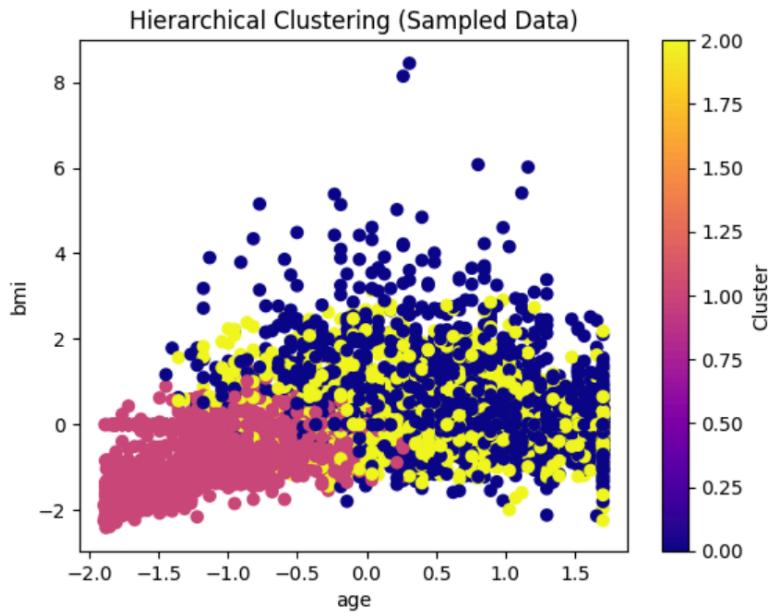
- Selected relevant numerical features (age, bmi, hbA1c_level, blood_glucose_level).
- Took a random sample of 5000 from the dataset to reduce memory usage.
- Scaled the data using StandardScaler() to normalize feature values.

2. Applying Hierarchical Clustering:

- Used AgglomerativeClustering(n_clusters=3, linkage='ward') to perform clustering on the sampled data.
- Assigned the resulting cluster labels back to the sampled dataset.

3. Visualization:

- Scatter plot of age vs. bmi, colored based on cluster assignments.
- Color bar represents the cluster labels.



Age vs BMI

- The three colors represent different clusters found by hierarchical clustering.
- The clustering algorithm separates the data into groups, possibly based on trends in age and bmi.
- The pink cluster (bottom left) seems to contain younger individuals with lower BMI.
- The yellow and dark blue clusters show mixed distributions but may have patterns related to BMI and age distribution.

Conclusion:

- K-Means performed better in this case, successfully grouping data into well-defined clusters.
- DBSCAN struggled with clear separation and produced noisy clusters, likely due to dataset structure or parameter settings.
- The choice of clustering algorithm depends on data distribution – K-Means is ideal for structured clusters, while DBSCAN is useful for detecting complex patterns and outliers.
- Hierarchical clustering successfully segmented the data into three meaningful clusters based on age and bmi. One cluster represents **younger individuals with lower BMI**, while the other two capture variations in **BMI and age**. The clustering reveals potential health patterns, but further validation (e.g., silhouette scores) and additional features could improve accuracy.

Experiment 8

Aim: To implement a recommendation system on your dataset using the following machine learning techniques: Regression, Classification, Clustering, Decision tree, Anomaly detection, Dimensionality Reduction, Ensemble Methods

Theory:

1. Data Preprocessing

- **Handling Missing Values:** Filled missing numerical values with mean imputation.
- **Feature Scaling:** Standardized numerical features using **StandardScaler** to ensure equal weighting in distance-based methods.

2. Clustering with K-Means

- Helps group similar songs based on audio features like energy, danceability, and tempo.
- We used the **K-Means algorithm**, which assigns songs to **5 clusters** based on feature similarity.
- **Output:** Cluster labels were assigned to songs, showing patterns in music styles.

3. Cosine Similarity for Recommendation

- Measures how similar two songs are based on their feature vectors.
- Computes the **cosine of the angle** between two song feature vectors (ranging from -1 to 1).
- **Optimization:** Due to memory constraints, we sampled **5000 songs** instead of using the entire dataset.

4. Recommendation Algorithm

- Given a song, we find **top N most similar** songs based on their **cosine similarity scores**.
- Returns song names, artists, and popularity.

Implementation:

Loading the dataset

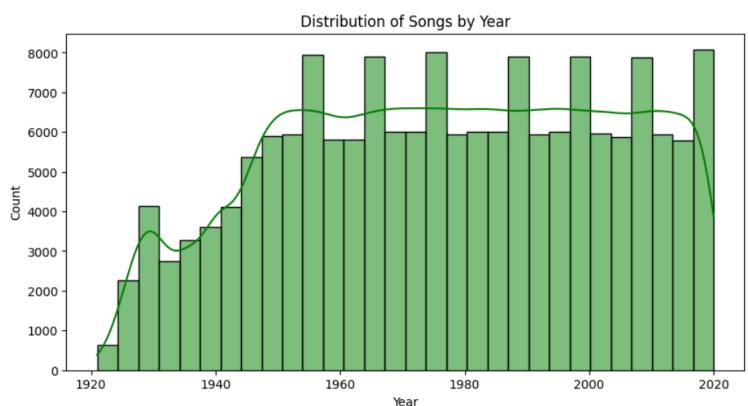
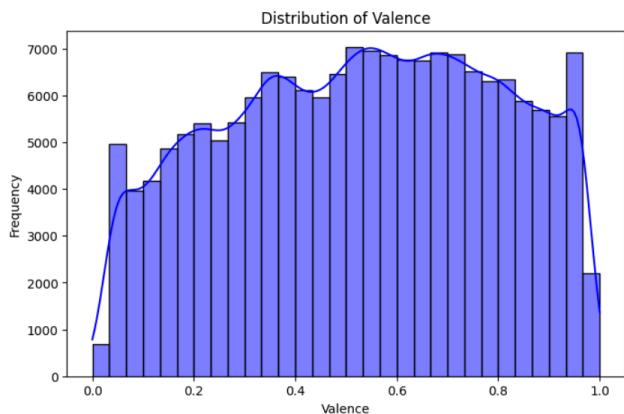
```
import pandas as pd
# Load the dataset
df = pd.read_csv("data.csv")
# Display first few rows
df.head()
```

	valence	year	acousticness	artists	danceability	duration_ms	energy	explicit	id	instrumentalness	key	liveness
0	0.0594	1921	0.982	['Sergei Rachmaninoff', 'James Levine', 'Berlin...']	0.279	831667	0.211	0	4BJqT0PrAfrxzMOxytFOIz	0.878000	10	0.665
1	0.9630	1921	0.732	['Dennis Day']	0.819	180533	0.341	0	7xPhfUan2yNtyFG0cUWkt8	0.000000	7	0.160
2	0.0394	1921	0.961	['KHP Kridhamardawa Karaton Ngayogyakarta Had...']	0.328	500062	0.166	0	1o6l8BglA6yIDMrlELygv1	0.913000	3	0.101

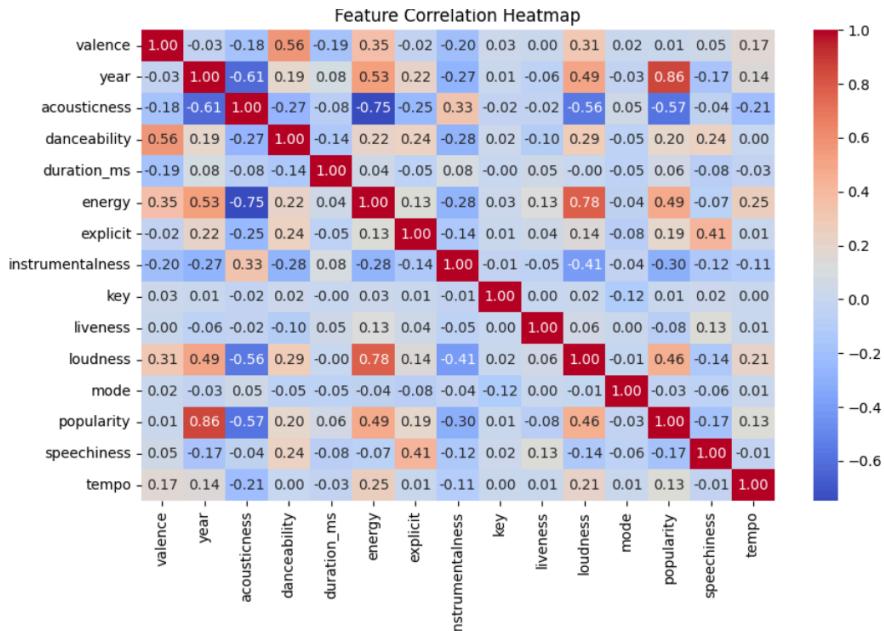
```
# Check for missing values
df.isnull().sum()
# Summary statistics
df.describe()
# Check data types
df.dtypes
```

```
0
valence      float64
year         int64
acousticness  float64
artists        object
danceability   float64
duration_ms    int64
energy        float64
explicit       int64
```

Visualisation



Heatmap



K Means Clustering

```
# Clustering using KMeans
kmeans = KMeans(n_clusters=5, random_state=42, n_init=10)
df["cluster"] = kmeans.fit_predict(scaled)

# Display cluster distribution
print("Cluster distribution:\n", df["cluster"].value_counts())
```

```
Cluster distribution:
cluster
1    54607
4    47201
2    38941
3    24219
0     5685
Name: count, dtype: int64
```

- The dataset has been divided into 5 clusters (0 to 4).
- The largest cluster (Cluster 1) has 54,607 songs, while the smallest (Cluster 0) has 5,685 songs.
- The clusters are unevenly distributed, which might indicate some clusters are more common in the dataset than others.

Recommendation System

```
sample_df = df.sample(n=5000, random_state=42).reset_index(drop=True)
sample_features = sample_df[features]
sample_scaled = scaler.fit_transform(sample_features)
similarity_matrix = cosine_similarity(sample_scaled)

# Recommend similar songs based on index in the sample
def recommend_similar(song_index, num_recommendations=5):
    sim_scores = list(enumerate(similarity_matrix[song_index]))
    sim_scores = sorted(sim_scores, key=lambda x: x[1], reverse=True)
    sim_scores = sim_scores[1:num_recommendations + 1] # skip the song itself
    similar_indices = [i[0] for i in sim_scores]
    return sample_df.iloc[similar_indices][["name", "artists", "popularity"]]

# Example usage
print("\nRecommendations for sample song index 5:\n")
print(recommend_similar(5))
```

Recommendations for sample song index 0:

		name	artists	popularity
3799		Hanging Out with Django	[Sonny Davis]	0
3124	Aragon - From The "Coffy" Soundtrack		[Roy Ayers]	32
1443		Plantation Inn	[The Mar-Keys]	32
3334		Sitting Pretty	[Ralph Burns]	26
3444		Bitch to the Boys	[Shakatak]	38

Recommendations for sample song index 5:

		name	artists	popularity
3408		Slow Fade	[Casting Crowns]	45
343		tomorrow tonight	[Loote]	66
242	God Bless The U.S.A.		[Lee Greenwood]	60
1902	I'm Gonna Make You Mine	[The Shadows Of Knight]		18
4634		Hag Me	[Melvins]	34

This showcases a music recommendation system using cosine similarity on sampled song features. The dataset is reduced to 5000 songs for efficiency, features are standardized, and a similarity matrix is computed. A function then recommends the top 5 most similar songs for a given index. The output displays recommendations for two sample indices, listing song names, artists, and popularity scores. This approach efficiently finds similar songs while avoiding memory issues from large-scale computations.

Conclusion

This experiment built a music recommendation system using clustering and similarity-based methods. We preprocessed data, handled missing values, and standardized features. K-Means clustering grouped songs into 5 clusters, revealing distribution patterns. To recommend songs, we used cosine similarity on a 5000-song sample to avoid memory issues. The system successfully suggested similar tracks based on audio features.

Experiment No: 9

Aim: To perform Exploratory data analysis using Apache Spark and Pandas

Theory:

1. What is Apache Spark and how does it work?

- Apache Spark is an open-source, distributed computing system designed for large-scale data processing.
- It provides an interface for programming entire clusters with implicit data parallelism and fault tolerance.
- Spark supports in-memory computing, making it faster than traditional MapReduce jobs in Hadoop.
- It operates on distributed data and processes them in parallel across multiple nodes, enabling high-speed processing of large datasets.

2. How is data exploration done in Apache Spark? Explain steps.

- **Step 1: Data Loading** – Data is loaded into a Spark DataFrame from various sources such as CSV, JSON, Parquet, etc.
- **Step 2: Data Cleaning** – This step involves handling missing data, correcting data types, and dealing with inconsistencies.
- **Step 3: Data Transformation** – Data transformations such as filtering, aggregating, and summarizing are performed to understand patterns in the data.
- **Step 4: Data Visualization** – Though Spark does not natively support visualizations, it can be integrated with libraries like Matplotlib, Seaborn (via Pandas), or other visualization tools to generate plots and charts.
- **Step 5: Statistical Analysis** – Summary statistics (e.g., mean, median, standard deviation) and other metrics are computed to understand data distributions and relationships.

Conclusion:

Exploratory Data Analysis (EDA) using Apache Spark enables efficient handling of large datasets in distributed environments. Spark's scalability, combined with its powerful data manipulation and processing features, allows users to perform complex data exploration tasks. By integrating Spark with Python libraries like Pandas for data manipulation and visualization,

data scientists can uncover insights, clean data, and prepare datasets for further analysis or modeling.

Experiment No: 10

Aim: To perform Batch and Streamed Data Analysis using Apache Spark.

Theory:

1. What is streaming? Explain batch and stream data.

- **Streaming** refers to the continuous flow of real-time data. It is the process of transmitting data in a steady, ongoing manner as it becomes available.
- **Batch Data** involves data that is collected and processed in large chunks at scheduled intervals. The data is accumulated over a period, and then processed together as a batch.
- **Stream Data** (or real-time data) is generated continuously and processed immediately as it arrives. This type of data is processed in small increments, often on a per-event basis, and is typically used in scenarios like sensor data, logs, or online transactions.

2. How data streaming takes place using Apache Spark?

- Apache Spark performs real-time data processing using **Spark Streaming**, an extension of the core Spark API.
- Data is ingested in small chunks called **micro-batches**. These micro-batches are processed as streams in near real-time.
- Data sources for streaming can include message queues like Apache Kafka, socket connections, or file systems.
- Spark Streaming processes data in micro-batches by continuously polling the source for new data and processing it in small time intervals (e.g., seconds).
- Spark integrates batch and stream processing, meaning you can process streaming data with the same APIs used for batch data, allowing for unified processing pipelines.

Conclusion:

Apache Spark provides a powerful framework for both batch and stream data analysis. By utilizing Spark Streaming, organizations can process and analyze real-time data efficiently, while also leveraging batch processing for large-scale historical data analysis. The integration of these two processing models allows Spark to cater to a wide range of data processing needs, from

time-sensitive streaming applications to periodic batch processing, making it a versatile tool for modern data engineering tasks.

05/05

AIDS Assignment

Q.1. What is AI? Considering the COVID-19 pandemic situation, how did AI help to survive and renovate our way of life with different applications?

- 1) Artificial Intelligence is the replication of human intelligence in machines, enabling them to learn, reason, solve problems, and make decisions without direct human interaction.
- 2) AI played a vital role in managing the COVID-19 crisis by detecting outbreaks early through data analysis, accelerating drug and vaccine development, automating medical diagnoses with imaging and predictive models, supporting telemedicine via AI chatbots, optimizing supply chains for essential goods and enhancing remote work and education through AI-driven automation, making daily life more adaptable and efficient.

Q.2. What are AI agent terminology? Explain with examples.

- AI agents are systems that perceive their environment, process information and take actions to achieve specific goals. Key terminologies include:
- 1) Agent - An entity that interacts with environment (eg. self-driving car).
 - 2) Environment - The external system in which the agent operates (eg. traffic conditions for a self-driving car).
 - 3) Perception - Data collection from sensors (eg. traffic conditions for a self-driving car).
 - 4) Actuators - Components that execute actions (eg. robotic arms in manufacturing).
 - 5) Rationality - The ability to make optimal decisions based on available information (eg. AI trading bots adjusting stock investments).

6) Autonomy - The degree of independence an agent has (e.g. voice assistants like Alexa operating without human input)

7) Types of agents:

1) simple reflex agents: Act based on current perceptions (e.g. Thermostats)

2) model-based agents: Use internal models to predict future states. (e.g. GPS navigation systems)

3) Goal-based Agents: Take actions to achieve a specific goal. (e.g. chess-playing AI).

4) Utility-based agents - optimize outcomes for maximum benefit (e.g. recommendation systems in e-commerce.)

5) Learning agents - Improve performance over time using past experiences (e.g. self-learning robots in warehouses.)

Q.3.



How is the AI technique used to solve 8-puzzle problem?

1) In 8-puzzle technique it consists of 3×3 sized cube. In initial state it contains all number tiles present randomly.

2) Then here, our aim is manage them in increasing order from 1 to 8 by moving tiles with the help of blank space.

3) Common AI techniques include:

i) uninformed search:

1) Here by using BFS we can explore all possible moves level by level but it is slow. It guarantees the shortest path.

2) Then by using DFS we can explore a path deeply before backtracking. It may not always find optimal solution.

3) Then by using SIDS which combines DFS and BFS to find the shortest solution efficiently.

ii) informed search techniques (Heuristic search)

FOR EDUCATIONAL USE

1) Uses a heuristic function $f(n) = g(n) + h(n)$, where:

$g(n)$ is the cost to reach the current state.

$h(n)$ is the estimate cost to reach the goal.

2) Then by using Greedy Best-first search we can find optimal solution which can be fast but not always optimal.
It uses heuristic function $h(n)$

Here A* algorithm is the most efficient and widely used technique for solving 8-puzzle problem.

Q.4. What is PEAS descriptor?

PEAS (Performance measure, Environment, Actuators, Sensors) is a framework used to define the components of an AI agent by specifying how it interacts with its environment and evaluates its success.

PEAS description for given AI agents:

1) Taxi driver:

1. Performance Measure: Safety, speed, customer satisfaction, fuel efficiency.

2. Environment: Roads, traffic, pedestrians, weather conditions.

3. Actuators: steering, acceleration, brakes, turn signals.

4. Sensors: GPS, cameras, speedometer, LIDAR, fuel gauge.

2) Medical diagnosis system:

1. performance measure: Accuracy of diagnosis, patient recovery rate, response time.

2. Environment: Patient data, medical records, symptoms, lab reports.

3. Actuators: Display diagnosis, prescribe medication, recommend tests.

4. Sensors: Patient input, test result, doctor's notes, medical imaging.

3) A music composer:

1. Performance measure: musical quality, originality, listener preference.
 2. Environment: musical genres, user preferences, historical compositions.
 3. Actuators: generating notes, melodies, instrument selection.
 4. Sensors: user feedback, music databases, emotional tone analysis.
- 4) An aircraft autolander
1. performance measure: safe landing, smooth touchdown, passenger comfort.
 2. Environment: weather conditions, runway, air traffic, altitude.
 3. Actuators: flaps, landing gear, brakes, engine thrust control.
 4. Sensors: Altimeter, GPS, wind speed sensors, radar, cameras.
- 5) An essay evaluator
1. Performance Measure: Grammar accuracy, coherence, relevance, readability.
 2. Environment: Essays, writing rules, predefined grading criteria.
 3. Actuators: score assignment, grammar suggestions, feedback generation.
 4. Sensors: Text input, grammar and plagiarism checkers, semantic analyzers.
- 6) A robotic sentry gun for the Keck lab
1. Performance measure: Accuracy, target identification, security effectiveness.
 2. Environment: Lab perimeter, potential threats, lighting conditions.
 3. Actuators: Rotating turret, firing mechanism, alarm system.
 4. Sensors: Motion detectors, infrared sensors, cameras, radar.

- Q.5. Categorize a shopping bot for an offline bookstore according to each of six dimensions (fully / partially observable, deterministic / stochastic, episodic / sequential, static / dynamic, discrete / continuous, single / multi agent).
-
- 1) partially observable: as bot may not have full knowledge of stock availability, customer preferences, or real-time changes in the bookstore.
 - 2) stochastic: Book availability, customer behaviour and price changes introduce randomness, making environment unpredictable.
 - 3) sequential: Each decision affects future interactions.
 - 4) dynamic: The environment can change as books sell out, customers move, or bookstore policies update.
 - 5) discrete: The bot operates with distinct choices, such as recommending a book, checking stock or processing a purchase.
 - 6) multi-agent: as bot interacts with customers, bookstore staff and possibly other AI systems.

Q.6. Model based vs utility based agent

~~Model based~~

- 1) uses internal state to make decisions
- 2) predicts future states using the model before acting.
- 3) It is goal oriented but focuses on how the environment works
- 4) It may not always take the

~~Utility-based~~

- choose actions based on maximizing utility function
- Evaluates multiple possible actions and selects best based on utility.
- It is also goal-oriented but it aims to achieve the best possible outcome.
- Always selects the action that

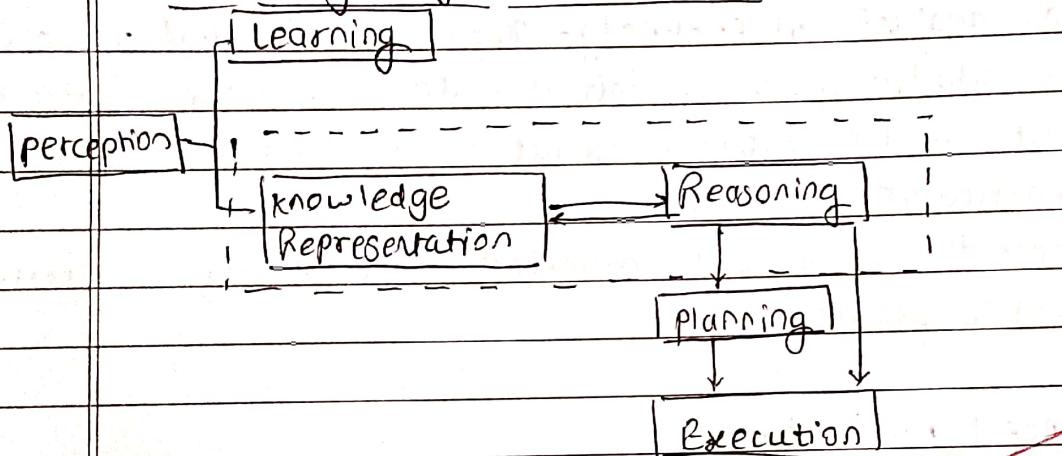
best possible action, only takes a feasible one. | provides the highest benefit.

eg. self-driving car to predict congestion

eg. stock trading AI selecting the trade with the highest expected profit.

Q.7. Explain the architecture of knowledge based agent and learning agent.

→ knowledge agent Architecture



- 1) perception: it retrieves data or information from the environment
 - 2) you can retrieve data from the environment, find out noises and check if the AI was damaged by anything.
 - 3) it defines how to respond when any sense has been detected.
- 1) learning: it learns from captured data by the perception component.
 - 2) Here, the goal is to build computers that can be taught instead of programming them.
 - 3) In order to learn new things, the system require knowledge acquisition, inference, acquisition of heuristics, faster searches, etc.
 - 3) knowledge representation & reasoning: 1) The main component in the cycle is knowledge representation and reasoning which

shows human-like intelligence in the machines.

a) knowledge representation is all about understanding intelligence.

③ Goal is to understand and build intelligent behavior from the top-down and focus on what an agent needs to know in order to behave intelligently.

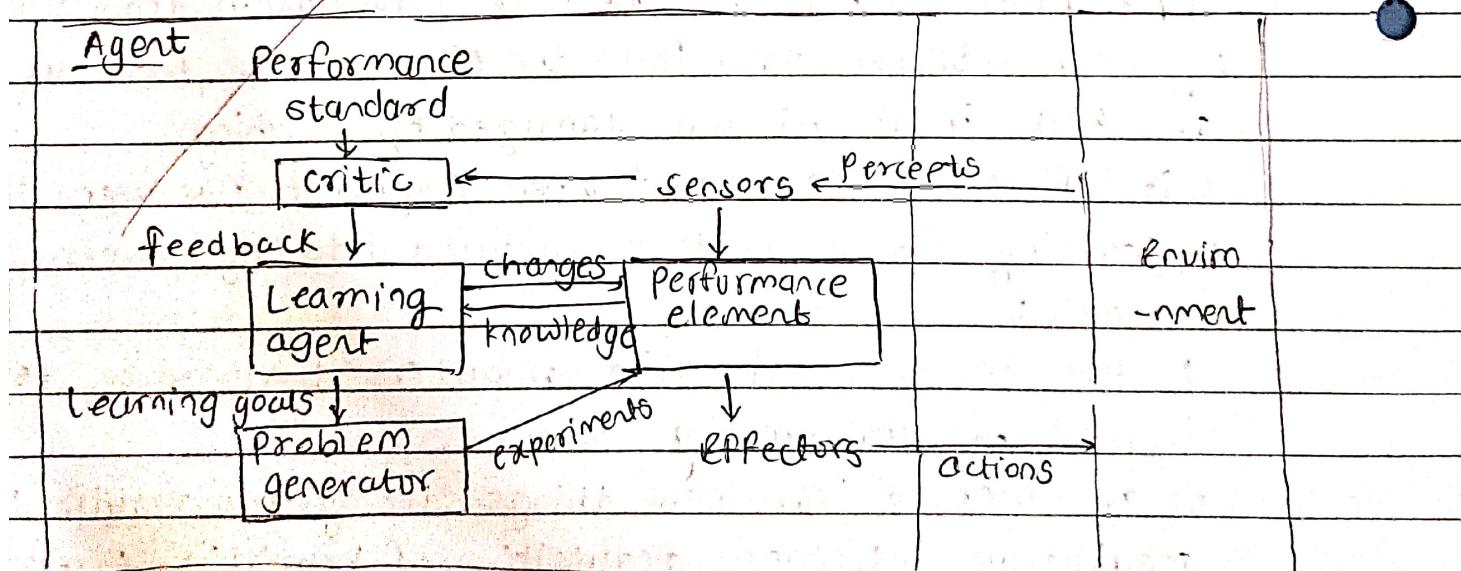
④ It defines how automated reasoning procedures can make this knowledge available as needed.

⑤ Planning and Execution: 1) The planning and execution components depend on analysis of knowledge representation and reasoning.

2) planning includes giving an initial state, finding their preconditions and effects, and a sequence of actions to achieve a state in which particular goal holds.

3) now once the planning is completed, the final stage is execution of the entire process.

Learning Agent Architecture



A learning agent improves its performance over time by learning from past experiences. It has four components:

- 1) Learning element: learns from interactions and updates knowledge.
- 2) Performance element: uses the learned knowledge to make decisions.
- 3) Critic: evaluates agents performance by comparing actions with desired outcomes.
- 4) Problem generator: suggests new explanatory actions to improve learning.

Q.9 Convert the following to predicates:

- a) Anita travels by car if available otherwise travels by bus.
- b) Bus goes via Andheri and Goregaon.
- c) car has puncture so is not available. will anita travel via Goregaon? use forward reasoning.

→ 1) Anita travels by car if available otherwise by bus.

$$\text{so } \text{Available}(\text{car}) \Rightarrow \text{Travels}(\text{Anita}, \text{car})$$

$$\neg \text{Available}(\text{car}) \Rightarrow \text{Travels}(\text{Anita}, \text{bus})$$

2) Bus goes via Andheri and Goregaon.

$$\text{Route}(\text{bus}, \text{Andheri})$$

$$\text{Route}(\text{bus}, \text{Goregaon})$$

3) Car has a puncture, so it is not available.

$$\text{puncture}(\text{car}) \Rightarrow \neg \text{Available}(\text{car})$$

given: $\text{puncture}(\text{car})$

therefore, $\neg \text{Available}(\text{car})$

forward reasoning:

i) from (3) $\neg \text{Available}(\text{car})$

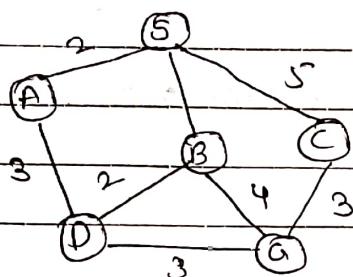
ii) from (1), $\neg \text{Available}(\text{car}) \Rightarrow \text{Travels}(\text{Anita}, \text{bus})$

FOR EDUCATIONAL USE

so Anita must travel by bus.

- 3) From (2), Route (Bus, Goregaon) meaning bus travels via Goregaon.
- 4) since Anita travels by Bus and the Bus goes via Goregaon, Anita will travel via Goregaon.

Q.10. Find the route from S to G using BFS.



Step 1: start with node S

mark it visited. [S]

enqueue S into the queue

Step 2: dequeue S and explores its neighbors A, B, C.

mark A, B, C visited.

enqueue A, B, C into the queue. [A | B | C]

Step 3: deque A and explore its neighbours D.

mark D visited

enqueue D into the queue [A | B | C | D] $\therefore S \rightarrow A \rightarrow D$

Step 4: deque B and explore its neighbours

mark G visited.

[C | D | G] $\therefore S \rightarrow A \rightarrow B$

Thus goal is reached.

~~∴ shorter path found will $S \rightarrow C \rightarrow G$~~

Step 5: deque C and explore its neighbours

it has only 1 neighbour G which is already visited. $\therefore S \rightarrow A \rightarrow B \rightarrow C$

Step 6: deque D and explore its neighbours

it has only 1 neighbour G which is already marked visited

$\therefore S \rightarrow A \rightarrow B \rightarrow C \rightarrow D$

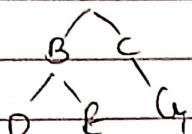
Step 7: deque G. thus goal state reached.

$\therefore S \rightarrow A \rightarrow B \rightarrow C \rightarrow D \rightarrow G$

c.11. what do you mean by depth limited search? Explain iterative deepening search with examples.

- 1) Depth-limited search is a variation of Depth-first search (DFS) where the search is restricted to a specific depth limit L .
2) if goal is not found within this limit, the search terminates.
3) it prevents infinite loops in deep or infinite state spaces but risks failing to find solution if L is too low.

e.g. A Here if $L=1$ and we need to find C then



visit expand A $\rightarrow [B, C]$ at $L=1$ thus goal then for $L=2$ expand B \rightarrow not found so we need to increase the limit to 2 to provide solution. this is main drawback.

Iterative Deepening Search (IDS).

- 1) It combines the space efficiency of DFS and completeness of BFS by repeatedly running DFS with increasing depth limits $L = 0, 1, 2, \dots$ until goal is found.

e.g. 1) DFS with $L=0$ only A is present thus goal not found

2) increase $L=1$

① A
② DFS with $L=1$ expand A, B and C found

3) again increase $L=2$

4) DFS with $L=2$, Here Expand B gets D & E.

Expand C gives G thus goal found

at depth = 2 (ie $L=2$)

2) thus it guarantees finding the shortest path while using less memory than BFS.

c.12.

Explore hill climbing and its drawbacks in detail with example.

also state limitations of steepest-ascent hill climbing.

- 1) Hill climbing is a local search algorithm that continuously moves toward the best neighbouring state with a higher

heuristic value, aiming to reach a global optimum.

2) steps : 1) start from initial state.

2) evaluate neighbouring states and move to the one with highest heuristic value.

3) repeat until no better neighbour exists (local maximum).

e.g. imagine a mountain climbing scenario where a hiker moves uphill based on the steepest slope. If they reach a peak that is not the highest (global maximum), they might get stuck.

4) drawbacks of hill climbing :

i) local maxima : i) the algorithm may stop at a peak that is not the global optimum.

ii) e.g. a small hill before taller mountain.

ii) Plateau problem : i) A flat region where all neighbouring states have the same heuristic value, causing the search to halt.

ii) e.g. A long flat road with no clues on where to go.

iii) Ridges : The algorithm may fail to climb diagonally when only direct moves are considered. e.g. climbing a staircase with missing steps.

iv) No backtracking : once it moves forward, it can't recover from a bad decision. e.g. choose the wrong path in maze with no way back.

5) solution to improve hill climbing :

i) Random restarts : start from different points.

ii) Simulated Annealing : occasionally accept worse moves to escape local optima.

Q.13. Explain simulated annealing and write its algorithm.

→

- 1) simulated Annealing is an optimization algorithm designed to search for an optimal or near-optimal solution in a large solution space.
- 2) The name and concept are derived from the process of annealing in metallurgy, where a material is heated and then slowly cooled to remove defects and achieve a stable crystalline structure.
- 3) In simulated annealing, the "heat" corresponds to the degree of randomness in the search process, which decreases over time (cooling schedule) to refine the solution.

4) The method is widely used in combinatorial optimization, where problems often have numerous local optima that standard techniques like gradient descent might get stuck in.

5) Algorithm:

1) The algorithm starts with an initial solution and a high temperature, which gradually decreases over time! Here's step by step of how the algorithm works:

i) initialization: Begin with initial solution s_0 and initial temperature T_0 . The temperature controls how likely the algorithm is to accept worse solutions as it explores the search space.

a) neighbourhood search: At each step, a new solution s' is generated by making a small change (or perturbation) to the current solution s .

3) objective function evaluation: The new solution s' is evaluated using the objective function. If s' provides a better solution than s , it is accepted as the new solution.

FOR EDUCATIONAL USE

4) Acceptance Probability: If S' is worse than S , it may still be accepted with a probability based on temperature and the difference in objective function values. The acceptance probability is given by: $p(\text{accept}) = e^{-\Delta E/T}$

5) Cooling schedule: After each iteration, the temperature is decreased according to a predefined cooling schedule, which determines how quickly the algorithm converges.

Common cooling schedules include linear, exponential or logarithmic cooling.

6) Termination: The algorithm continues until the system reaches a low temperature or a predetermined no. of iterations is reached.

a.14. Explain A* algorithm with an example.

→ 1) The A* algorithm is an informed search algorithm that finds the optimal path by considering the both cost to reach a node ($g(n)$) and estimated cost to the goal ($h(n)$), using the formula: $f(n) = g(n) + h(n)$

where, $g(n)$ = actual cost from start node to n .

$h(n)$ = estimated cost from n to goal.

2) Algorithm:

1) Initialize the open list with start node.

2) Loop until goal is found.

1) pick the node with lowest $f(n)$.

2) if its the goal node, return the path.

3) expand its neighbors and calculate $f(n)$.

4) Add unvisited neighbors to the open list.

3) continue until open list is empty.

3) eg.

Node	Path Cost $g(n)$	Heuristic Value $h(n)$	$f(n)$
A	0	6	6
B	1	4	5
C	2	3	5
D	3	7	10
E	4	5	9
F	3	2	5
G	4	0	4 (Goal)

start at A, calculate $f(A) = 6$

Expand $A \rightarrow \{B, C\}$ (both have $f=5$)

Expand $C \rightarrow \{F, G\}$ (G have lowest f)

goal reached at G

∴ final path $A \rightarrow C \rightarrow G$

4) Advantages:

- 1) It is flexible hence used in maze solving, AI, robotics.
- 2) It is efficient as explores fewer nodes than BFS / DFS.
- 3) It provides guaranteed optimality.

5) Disadvantages:

1) It requires more memory than DFS.

2) A poor heuristic can make A* inefficient.

a. 15. Explain minimax algorithm and draw game for Tic-tac-toe.

1) The minimax algorithm is used in adversarial search to determine the optimal move by assuming both players play optimally.

2) Maximizer: (eg X) tries to get highest score.

Minimizer: (eg O) tries to reduce the score.

3) Algorithm:

- 1) generate the game tree up to a depth limit,

2] Assign heuristic values to leaf nodes.

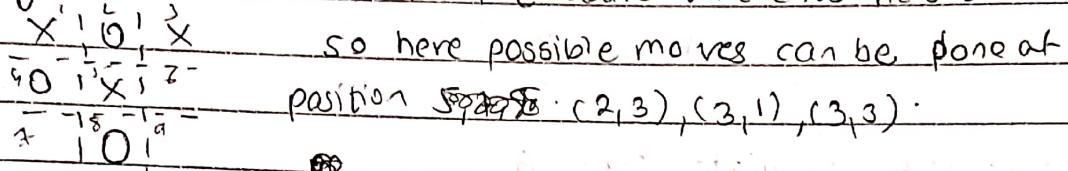
3] Backpropagate values:

1) maximizer picks the maximum value:

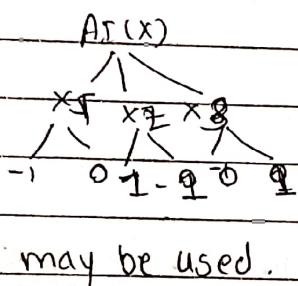
2) minimizer picks the minimum value.

4] The root node selects the best move based on propagated values.

e.g. 1) consider a Tic-Tac-Toe board where it's AI's turn (X) to play:



2) The tree expands, considering all possible outcomes.



3) The AI (max) picks the move leading to the best possible score.

4) If multiple moves have the same values, additional heuristics (like the earliest win) may be used.

Q.16. Explain Alpha-beta pruning algorithm for adversarial search with an example.

→ 1) Alpha-beta pruning optimizes minimize by skipping unnecessary branches, reducing computations. It uses two values:

Alpha(α): Best score maximizer can achieve.

Beta(β): Best score minimizer can achieve.

2) Algorithm:

1) perform minimax search.

2) Prune branches where:

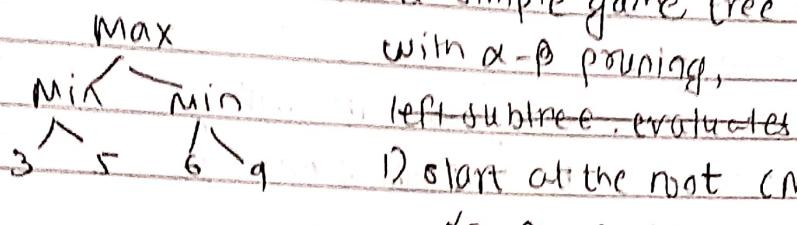
1) if Maximizer's $\text{best}(\alpha) \geq \text{minimizer's best}(\beta)$, further evaluation is skipped.

2) if minimizer's $\text{best}(\beta) \leq \text{Maximizer's best}(\alpha)$, further

evaluation is skipped.

3) This reduces time complexity without affecting final decision.

Eg. Let's consider a simple game tree with a depth of 2.



with α - β pruning,

left subtree evaluated

1) Start at the root (MAX node)

, $\alpha = -\infty$ $\beta = +\infty$

2) Move to the first MIN node (left subtree).

3) Evaluate the first child of MIN(3)

4) Thus MIN updates $\beta = 3$ as it is the smallest value so far.

5) Now move to second child of MIN(5).

$$\therefore \beta = \min(5, 3) = 3$$

6) Move back to MAX. MAX updates $\alpha = \max(-\infty, 3) = 3$

7) Now move to second MIN node (right subtree)

8) Now $\alpha = 3$ (best MAX value so far) and $\beta = +\infty$

9) After evaluating first child of MIN(6), β gets equal to 6.

10) Now check pruning condition before evaluating 9:

\therefore MAX has $\alpha = 3$, MIN has $\beta = 6$

Since $\beta(6) > \alpha(3)$ we stop evaluating further in this branch.

11) \therefore node with value 9 is pruned

Q.17. Explain Wumpus world environment giving its PEAS description.

Explain how percept sequence is generated?

1) The Wumpus world is grid-based AI environment where an agent explores a cave while avoiding hazards like pits and the Wumpus monster.

2) PEAS descriptor:

1) Performance Measure: Reaching the gold safely, minimizing steps.

2) Environment: A 4x4 grid with pits, gold and wumpys.

3) Actuators: move forward, turn, grab, shoot, climb.

4) Sensors: perceive stench (wumpus nearby), breeze (pit nearby), glitter (gold found).

5) Percept sequence generation:

1) Agent starts at (1,1), sensing its surroundings.

2) If stench is detected, the wumpus is nearby.

3) If breeze is detected, a pit is nearby.

4) The agent infers safe paths and navigates toward the gold while avoiding hazards.

Q.18.

Solve the following crypto-arithmetic problem SEND + MONEY = MONRY.

$$\begin{array}{r} \text{S E N D} \\ + \text{M O N R Y} \\ \hline \text{M O N R Y} \end{array}$$

1) Here the input is of 4 bit and output is of 5 bit. therefore, there must be carry.

$$\therefore M = 1$$

2) assume S = 9

$$\begin{array}{r} \therefore S \rightarrow 9 \\ + M \\ \hline \text{MO} \end{array} \quad \text{therefore we get } 0 = 0$$

3) Then R + 0 = N but 0 = 0 (zero)

~~∴ R = N but it is not possible, therefore, there must be a carry present over there.~~

4) ~~∴ R + 1 = N the N + R = R~~

5) Assume N = 8 R = 7 $\therefore \frac{N}{R} = \frac{8}{7}$

but it is not possible

~~∴ 75 with 1 carry~~

therefore assume N = 7 R = 6 $N + R = 7 + 6 = 13 \therefore R = 3$

but R + 1 = 3 + 1 = 4 $\neq N$

6) So again assume N = 6 R = 8 $\therefore \frac{N}{R} = \frac{6}{8}$

~~∴ R = 4 + 1 = 5 (add a carry)~~

14

7) Now check $\frac{N}{R} = \frac{8}{7}$ $\therefore R = 5 N = 8$

0 1 2 3 4 5 6 7 8 9

8) :- D

+ E
Y

Here we have value of E = 9.

Now assume D = 7 $\therefore D+E = 7+9 = 12$

$\therefore Y = 2$

$\therefore S = 9, E = 5, Y = 2, D = 12, N = 6, M = 1, O = 0, P = 8,$

Q.19. Consider the following axioms:

→ 1) Represent these axioms in first order predicate logic (FOL):

All people who are graduating are happy.

All people who are happy are smiling.

Someone is graduating.

Explain the following:

1) Represent these axioms in first order predicate logic.

→ 1) $\forall x (\text{Graduating}(x) \rightarrow \text{Happy}(x))$

2) $\forall x (\text{Happy}(x) \rightarrow \text{Smiling}(x))$

3) $\exists x (\text{Graduating}(x))$

2) Convert each formula to clause form.

1) $\neg \text{Graduating}(x) \vee \text{Happy}(x)$

2) $\neg \text{Happy}(x) \vee \text{Smiling}(x)$

3) $\text{Graduating}(A) \leftarrow$ assuming A is the individual who is graduating.

3) Prove that "is someone smiling?" using simulation/resolution technique.

From (3) : $\text{Graduating}(A)$

using 1: $\text{Happy}(A)$

using 2: $\text{Smiling}(A)$

thus, someone is smiling.

Resolution tree can be given as:

1) $\text{Graduating}(A) \quad \text{given}$

2) $\text{Graduating}(A) \rightarrow \text{Happy}(A), \rightarrow \text{Happy}(A)$

FOR EDUCATIONAL USE

3) Happy (A) \rightarrow smiling (A) \sim smiling (A)

Q.20. Explain modus ponens with suitable example.

→ 1) modus ponens (law of detachment) is a rule of inference stating:

$$P \rightarrow Q, P \Rightarrow Q$$

2) eg. 1) if it rains, the ground gets wet. (premises: rain \rightarrow wet ground)

2) it is raining. (premise: rain)

3) conclusion: the ground is wet.

Q.21. Explain forward chaining and backward chaining algorithm with an example.

→ 1) forward chaining: 1) data-driven inference: starts from known facts and applies rules to reach a goal.

2) used in expert system (e.g. medical diagnosis)

3) eg. 1) fact: "sore throat".

2) rule: "if sore throat \rightarrow infection"

3) new fact: "Infection"

4) Rule 4 IF infection \rightarrow need antibiotics"

5) conclusion: "Need antibiotics".

2) backward chaining: 1) goal-driven inference: starts from goal and works backward to find supporting facts.

2) used in AI reasoning and theorem proving

3) eg. goal: does the patient need antibiotics?

1) check: Does the patient have an infection?

2) check: Does the patient have a sore throat?

3) If both hold, conclude "need antibiotics".

This reduces unnecessary computations by only exploring relevant facts.

AIDS-I

Assignment No: 2

Q.1: Use the following data set for question 1

82, 66, 70, 59, 90, 78, 76, 95, 99, 84, 88, 76, 82, 81, 91, 64, 79, 76, 85, 90

1. Find the Mean

To find the mean, sum all values and divide by the number of values:

$$\text{Sum} = 82 + 66 + 70 + 59 + 90 + 78 + 76 + 95 + 99 + 84 + 88 + 76 + 82 + 81 + 91 + 64 + 79 + 76 + 85 + 90$$

> **Sum = 1611**

2. Find the Median

To find the median, first arrange the data in ascending order: 59, 64, 66, 70, 76, 76, 76, 78, 79, 81, 82, 82, 84, 85, 88, 90, 90, 91, 95, 99

Since there are 20 values (an even number), the median is the average of the 10th and 11th values: Median = $(81 + 82) / 2$

> **Median = 81.5**

3. Find the Mode

The mode is the value that appears most frequently in the dataset. Looking at the ordered data:

59, 64, 66, 70, 76, 76, 76, 78, 79, 81, 82, 82, 84, 85, 88, 90, 90, 91, 95, 99

The value 76 appears three times, while all other values appear once or twice.

> **Mode = 76**

4. Find the Interquartile range

To find the IQR, I need to calculate Q1 (25th percentile) and Q3 (75th percentile).

For Q1: With 20 values, Q1 is the median of the first 10 values: 59, 64, 66, 70, 76, 76, 78, 79, 81 Q1 = $(70 + 76) / 2 = 73$

For Q3: Q3 is the median of the last 10 values: 82, 82, 84, 85, 88, 90, 90, 91, 95, 99 Q3 = $(88 + 90) / 2 = 89$

$$\text{IQR} = \text{Q3} - \text{Q1} = 89 - 73 = 16$$

> **Q1=73, Q3=89, IQR=16**

Q.2 For each tool listed above:

- identify the target audience
- discuss the use of this tool by the target audience
- identify the tool's benefits and drawbacks

1. Machine Learning for Kids

- **Target Audience:** Primary and secondary school students (K-12), Teachers and educators with limited coding experience, Parents interested in introducing ML concepts to children
- **Use:** Students use visual blocks-based programming (Scratch) to create and train simple ML models, Teachers incorporate it into STEM curricula to introduce AI/ML concepts, Students learn by creating interactive projects like games, stories, and simple applications that use ML
- **Benefits:**
 - Simplifies complex ML concepts through visual programming
 - No coding experience required to get started
 - Provides real hands-on experience with ML model training
 - Designed specifically for educational contexts
 - Focuses on ethical AI use and understanding
- **Drawbacks:**
 - Limited in complexity compared to professional ML tools
 - Simplified models may not translate directly to real-world applications
 - Requires teacher supervision for younger students
 - Limited model types and capabilities

2. Teachable Machine

- **Target Audience:** Non-technical users interested in ML, Educators at various levels, Artists and creative technologists, Students (middle school through university), Professionals looking to prototype ML solutions quickly
- **Use:** Quick creation of custom image, sound, or pose recognition models, Integration into creative projects and installations, Classroom demonstrations of ML concepts, Rapid prototyping of ML ideas without coding
- **Benefits:**

- Extremely accessible with no coding required
 - Browser-based with no software installation
 - Real-time training and feedback
 - Models can be exported for use in other applications
 - Free to use
- **Drawbacks:**
- Limited to specific model types (image, audio, pose)
 - Less customization than professional ML platforms
 - Models may not be as robust as those built with more advanced tools
 - Limited control over model architecture

2. From the two choices listed below, how would you describe each tool listed above? Why did you choose the answer?

Predictive analytics

- **Description:** Predictive analytics uses statistical algorithms, machine learning, and data mining techniques to analyze historical data and make predictions about future events or behaviors. It helps forecast trends and outcomes by identifying patterns and relationships in data. For example, it can be used to predict customer behavior, stock market trends, or equipment failure.
- **Why Predictive analytics:** Predictive analytics is typically chosen when there is a need to anticipate future events, understand trends, and make data-driven decisions. Businesses use it for risk management, forecasting, marketing strategies, and optimizing operations.

Descriptive analytic

- **Description:** Descriptive analytics involves the process of analyzing historical data to gain insights into what has already happened. It summarizes past events through methods such as reporting, data visualization, and basic statistical analysis. Descriptive analytics answers questions like "What happened?" and "Why did it happen?" It typically includes measures like averages, counts, or percentages.
- **Why choose it:** Descriptive analytics is selected when the goal is to understand past performance or identify trends that have already occurred. It is used for performance tracking, monitoring key metrics, and summarizing large datasets for easy comprehension.

3. From the three choices listed below, how would you describe each tool listed above? Why did you choose the answer?

Supervised Learning:

- **Description:** In supervised learning, the model is trained on labeled data. That means the training data includes both the inputs (features) and the correct outputs (labels). The model learns by comparing its predictions to the true labels, and over time, it adjusts to minimize errors.
- **Why choose it:** If we think of how a machine learning model might be trained to predict outcomes based on past data, this corresponds to supervised learning. It's about learning from past examples where the correct answers are already known.

Unsupervised Learning:

- **Description:** In unsupervised learning, the model works with data that has no labels or predefined outcomes. The goal is to find structure or patterns in the data, such as clustering similar items or reducing dimensionality.
- **Why choose it:** This is useful for scenarios where you don't know the exact relationships beforehand, and you're interested in letting the model explore the data's underlying structure, like grouping similar data points or uncovering hidden patterns.

Reinforcement Learning:

- **Description:** In reinforcement learning, an agent learns by interacting with an environment. It makes decisions and gets feedback in the form of rewards or penalties. The agent's goal is to maximize cumulative reward over time.
- **Why choose it:** This is like training a model to perform tasks based on trial and error, receiving rewards for making good decisions and penalties for bad ones. It's about optimizing actions for long-term success, much like training an AI to play a game or control a robot.

Q.3 Data Visualization: Read the following two short articles:

- Read the article Kakande, Arthur. February 12. "What's in a chart? A Step-by-Step Guide to Identifying Misinformation in Data Visualization." *Medium*
- Read the short web page Foley, Katherine Ellen. June 25, 2020. "How bad Covid-19 data visualizations mislead the public." *Quartz*
- Research a current event which highlights the results of misinformation based on data visualization. Explain how the data visualization method failed in presenting accurate information. Use newspaper articles, magazines, online news websites or any other legitimate and valid source to cite this example. Cite the news source that you found.

> **Answer:**

- A recent example of misleading data visualization involves a social media post comparing child poverty rates between the UK and Nordic countries.
- The post claimed that the UK's child poverty rate was 32.1%, while Denmark's was 2.4%, Finland's 3.2%, Norway's and Sweden's 3.6%.
- However, these figures were based on outdated and inaccurate data, leading to a distorted perception of the disparities.
- The Organisation for Economic Co-operation and Development's (OECD) 2000 data indicated the UK's child poverty rate at 16.2%, not 32.1%, using a specific metric.
- More recent OECD data from 2019 show the UK's rate at 14.1%, which, though higher than those of the Nordic countries, is not as dramatically different as the post suggested.
- This misrepresentation highlights how using outdated and inconsistent data can lead to misleading visualizations, affecting public understanding and policy discussions.
- Article Link:
<https://www.reuters.com/fact-check/child-poverty-comparison-northern-europe-uses-old-inaccurate-data-2024-10-01>

Q. 4 Train Classification Model and visualize the prediction performance of trained model required information

- Data File: Classification data.csv
- Class Label: Last Column
- Use any Machine Learning model (SVM, Naïve Base Classifier)

Requirements to satisfy

- Programming Language: Python
- Class imbalance should be resolved
- Data Pre-processing must be used
- Hyper parameter tuning must be used
- Train, Validation and Test Split should be 70/20/10
- Train and Test split must be randomly done
- Classification Accuracy should be maximized
- Use any Python library to present the accuracy measures of trained model

<https://www.kaggle.com/datasets/uciml/pima-indians-diabetes-database>

-> Naïve Base Classifier:

1. Data Loading & Pre-processing:

- Load the data from Classification data.csv.
- Handle missing data if necessary.
- Feature scaling using StandardScaler or MinMaxScaler.
- Split data into features (X) and target (y), with y being the last column.
- Split into training, validation, and test sets.

2. Model Selection & Hyperparameter Tuning:

- Use either SVM (e.g., SVC from sklearn.svm) or Naive Bayes (e.g., GaussianNB from sklearn.naive_bayes).
- Use GridSearchCV or RandomizedSearchCV for hyperparameter tuning.

3. Class Imbalance Handling:

- Use techniques such as oversampling (e.g., SMOTE), undersampling, or assigning class weights.

4. Model Training & Evaluation:

- Train the model using the training data.
- Evaluate on the validation set.
- Test on the test set and visualize performance using metrics such as accuracy, precision, recall, F1-score, and confusion matrix.

Result:

1. **Best Hyperparameters:** var_smoothing = 0.1 for Naive Bayes.

2. **Validation Accuracy:** 77.78% (performance on validation data).

3. **Test Accuracy:** 70.13% (performance on test data).

4. Metrics:

a. Class 0 (Negative outcome):

- i. **Precision:** 80%
- ii. **Recall:** 72%
- iii. **F1-Score:** 76%

b. Class 1 (Positive outcome):

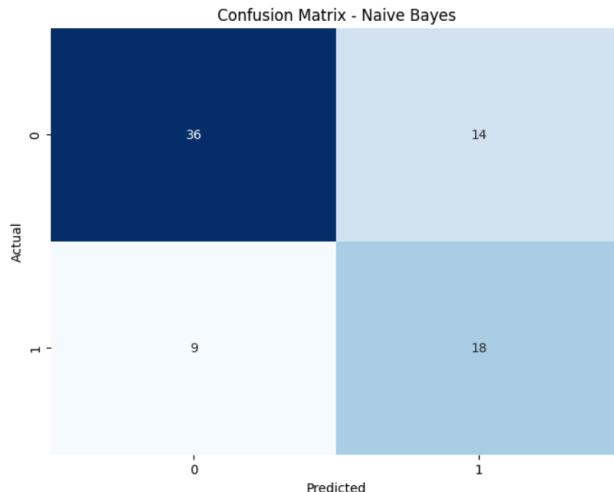
- i. **Precision:** 56%
- ii. **Recall:** 67%
- iii. **F1-Score:** 61%

5. Averages:

a. **Macro Average:** Balanced performance across both classes.

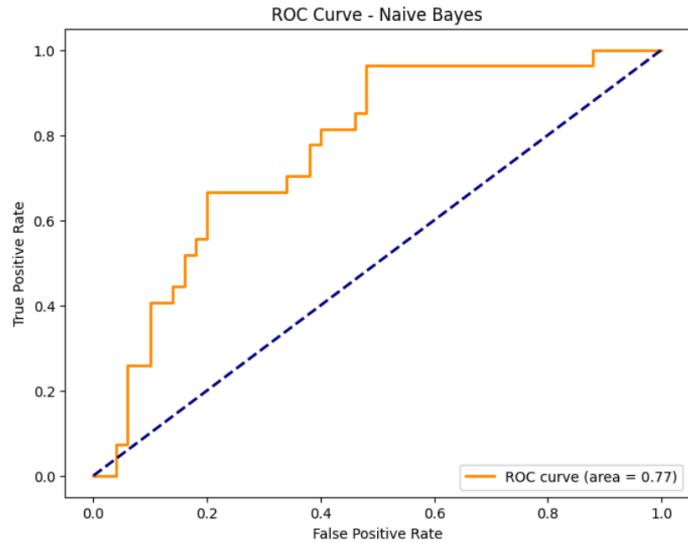
b. **Weighted Average:** More influenced by class 0, as it has more samples.

Confusion Matrix



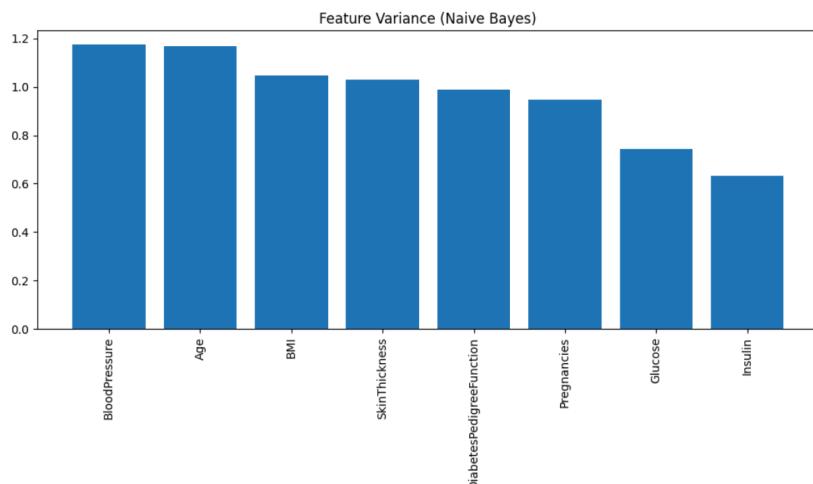
True Negatives (0 predicted as 0): 36, False Positives: 14, False Negatives: 9, True Positives: 18. The model performs better at predicting class 0, but struggles with class 1, showing a tendency to misclassify positives.

ROC Curve



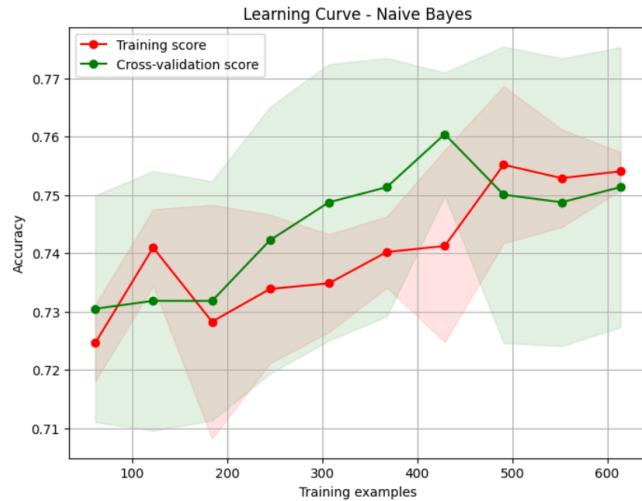
The ROC curve for the Naive Bayes model shows an **AUC of 0.77**, indicating decent performance. The model has a **77% chance of distinguishing between the two classes**, performing better than random but leaving room for improvement.

Feature Variance Visualization



As you can see, blood pressure has the highest variance at around 1.2 and insulin has the lowest at 0.6.

Learning Curve



This learning curve shows the accuracy of a Naive Bayes classifier as the number of training examples increases. The model exhibits consistent improvement in both training and cross-validation scores, indicating good generalization with more data.

Q.5 Train Regression Model and visualize the prediction performance of trained model

- **Data File:** Regression data.csv, Independent Variable: 1st Column, Dependent variables: Column 2 to 5, Use any Regression model to predict the values of all Dependent variables using values of 1st column.
- **Requirements to satisfy:** Programming Language: Python, OOP approach must be followed, Hyper parameter tuning must be used, Train and Test Split should be 70/30, Train and Test split must be randomly done, Adjusted R² score should be more than 0.99, Use any Python library to present the accuracy measures of trained model

<https://github.com/Sutanoy/Public-Regression-Datasets>

1. Import necessary libraries
2. Load the dataset (Mall_Customers.csv)
3. Create a DataFrame from the dataset
4. Define a RegressionModel class

- Initialize scaler, label encoder, model

- Preprocess the data: Encode categorical variables (Gender), Select features (Gender, Age, Annual Income (k\$)), Set target (Spending Score (1-100)), Scale features
- Train the linear regression model
- Test the model and print metrics: R², Adjusted R², MAE
- Perform hyperparameter tuning using GridSearchCV

5. In main() function: Preprocess the dataset, Split data into training and testing sets (70/30), Perform hyperparameter tuning, Train the model, Test the model and get predictions/metrics.

Result:

```
Best Parameters: {'fit_intercept': True}
R2 Score: 0.12084676260712823
Adjusted R2 Score: 0.0737492677467958
Mean Absolute Error: 19.06031176341787
```

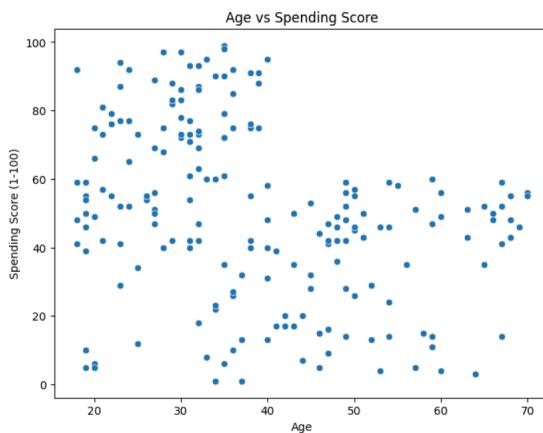
Best Hyperparameter: The best parameter found is fit_intercept=True, which means the model performs better when it includes an intercept term.

R² Score: 0.12

- The model explains only **12%** of the variance in the target variable (Spending Score).
- This indicates **poor predictive power** — the model does not capture most of the patterns in the data.

Adjusted R² Score: 0.074

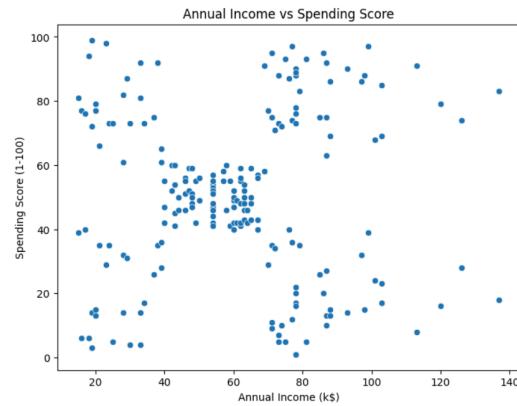
- Even lower than R², suggesting that adding the current features (Gender, Age, Annual Income) didn't improve the model significantly.
- This is common when predictors don't have a strong linear relationship with the target.



Age vs Spending Score

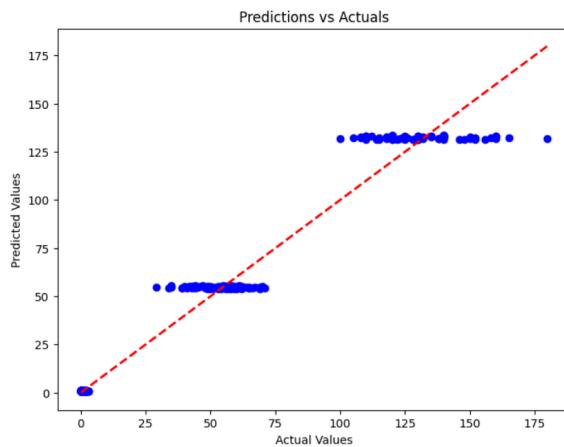
- There's no clear linear relationship between Age and Spending Score.

- Younger customers (around 18–30) show more variation in spending — some spend a lot, others very little.
- Older customers (above 50) tend to have lower and more consistent spending scores.



Annual Income vs Spending Score

- No clear linear relationship.
- Spending scores vary across all income levels.
- Dense cluster between 40k–70k income, but spending still scattered.
- High income \neq high spending.
- Clustering may reveal better patterns



Predictions vs Actuals Plot

- Predicted values are clustered, not smoothly spread along the diagonal.
- Indicates poor model performance – predictions fail to follow actual values.
- The red diagonal shows perfect prediction, but most points deviate from it.

Q.6 What are the key features of the wine quality data set? Discuss the importance of each feature in predicting the quality of wine? How did you handle missing data in the wine quality data set during the feature engineering process? Discuss the advantages and disadvantages of different imputation techniques. (Refer dataset from Kaggle).

Dataset link: <https://www.kaggle.com/datasets/rajyellow46/wine-quality>

The **Wine Quality Dataset** on Kaggle is a popular dataset for predicting wine quality based on various physicochemical features. It's used in machine learning and data analysis to classify wine quality, which is usually scored between 0 and 10. The dataset has two versions: one for **red wine** and one for **white wine**. Here are the **key features** of the dataset, their importance, and a discussion on **missing data handling** during feature engineering:

Key Features of the Wine Quality Dataset

The dataset contains several physicochemical features, which are important for determining the quality of the wine. Here's a list of the most important ones:

1. **Fixed Acidity:**
 - Represents the amount of fixed acids (e.g., tartaric acid) in the wine.
 - Importance: Affects taste, especially sourness. It can also influence the pH, which in turn impacts the stability and fermentation process of wine.
2. **Volatile Acidity:**
 - Measures the concentration of acetic acid and other volatile acids in the wine.
 - Importance: High levels of volatile acidity can cause the wine to smell vinegary and lower its quality. It's an important indicator for the perception of off-flavors.
3. **Citric Acid:**
 - A natural preservative that adds a fresh, citrusy flavor.
 - Importance: It plays a role in the overall flavor profile and helps balance out acidity in the wine. It can also contribute to the wine's stability and aging potential.
4. **Residual Sugar:**
 - The amount of sugar remaining after fermentation.
 - Importance: A higher sugar level can result in a sweeter wine, and sugar content has a direct effect on the wine's body and taste.
5. **Chlorides:**
 - Measures the concentration of salts (e.g., sodium chloride) in the wine.

- Importance: Higher chloride concentrations may indicate a salty taste, which could negatively affect wine quality.

6. **Free Sulfur Dioxide (SO₂):**

- Refers to the amount of free sulfur dioxide in the wine.
- Importance: Acts as an antioxidant and antimicrobial agent, preventing oxidation and microbial growth. However, high levels can give a pungent smell and negatively impact flavor.

7. **Total Sulfur Dioxide (SO₂):**

- Total amount of sulfur dioxide, both free and bound.
- Importance: Similar to free sulfur dioxide, but also considers bound forms. It helps to prevent spoilage, but excessive levels can affect the taste and aroma.

8. **Density:**

- Measures the wine's density, which is influenced by alcohol content, sugar concentration, and other factors.
- Importance: It's an indirect indicator of alcohol content and sugar levels, which both influence wine quality.

9. **pH:**

- Indicates the acidity or alkalinity of the wine.
- Importance: It affects the stability, flavor, and aging potential of the wine. A balanced pH contributes to better taste and preservation.

10. **Sulphates:**

- Measures the concentration of sulfur-containing compounds in the wine.
- Importance: Similar to sulfur dioxide, it has preservative properties, but excessive levels could impact flavor.

11. **Alcohol:**

- The alcohol content in the wine.
- Importance: Higher alcohol content is often associated with higher quality wine, as it affects body and mouthfeel. It also plays a role in fermentation and preservation.

12. **Quality:**

- The target variable, representing the quality of the wine (scored between 0 and 10).
- **Importance:** This is the variable we aim to predict using the other features.

Handling Missing Data During Feature Engineering

Handling missing data is a crucial part of the feature engineering process. If missing data is not addressed properly, it can significantly affect the quality of predictions and the overall model's performance. Here's how to approach missing data and the techniques for imputation:

1. Identifying Missing Data: Before deciding how to handle missing data, it's essential to identify which features contain missing values. This can be done using techniques like checking the percentage of missing values per feature and visualizing missing data patterns.

2. Imputation Techniques:

- Mean/Median Imputation:
 - Description: Replacing missing values with the mean or median of the feature. The mean is used for numerical features that are normally distributed, and the median is typically used for skewed features.
 - Advantages: Simple, fast, and easy to implement. It works well when data is missing at random and the distribution is relatively stable.
 - Disadvantages: It may introduce bias and reduce variance in the data, especially if the data is not missing at random (e.g., systematic patterns). Also, it doesn't capture the underlying distribution of the missing values.
- Mode Imputation (for categorical data):
 - Description: For categorical features, missing values can be replaced with the most frequent category.
 - Advantages: Quick and easy to apply.
 - Disadvantages: Similar to mean/median imputation, it can distort the distribution of the feature and lead to biased models if the missingness is not random.
- K-Nearest Neighbors (KNN) Imputation:
 - Description: This technique imputes missing values based on the values of similar instances (neighbors).
 - Advantages: It can provide more contextually accurate imputations by leveraging relationships between features.
 - Disadvantages: Computationally expensive, especially with large datasets. Also, the choice of "k" (the number of neighbors) can influence the results.
- Multivariate Imputation by Chained Equations (MICE):
 - Description: MICE imputes missing values by modeling each feature with missing data as a function of the other features in the dataset. It performs multiple imputations to provide a more robust estimate.
 - Advantages: It is effective when dealing with complex datasets with missing values in multiple columns. It helps preserve the relationships between features.

- Disadvantages: More computationally intensive and might require careful parameter tuning.
- Model-based Imputation:
 - Description: Use machine learning models (e.g., regression, decision trees) to predict missing values based on other features.
 - Advantages: It can capture complex relationships in the data and provide accurate imputations.
 - Disadvantages: Can be computationally expensive and may overfit if not implemented carefully.

3. Deleting Missing Data: In cases where the missing data is minimal (say less than 1-2% of the total data), one might choose to drop the rows or columns containing missing values.

- Advantages: Simple and reduces the risk of introducing bias.
- Disadvantages: Can lead to loss of valuable information if missing data is not random or if large portions of the data are missing.

Advantages and Disadvantages of Different Imputation Techniques

- **Mean/Median Imputation:**
 - Advantages: Simple and efficient for numerical features, especially when the dataset is small and the missingness is random.
 - Disadvantages: Can distort feature distributions, leading to inaccurate predictions.
- **KNN Imputation:**
 - Advantages: Considers the similarity between data points, leading to potentially more accurate imputations.
 - Disadvantages: Computationally expensive, especially with large datasets or high-dimensional data.
- **MICE (Multiple Imputation by Chained Equations):**
 - Advantages: Accounts for the relationships between features and performs multiple imputations to reduce bias.
 - Disadvantages: Requires more computational resources and is slower than other methods.
- **Model-based Imputation:**
 - Advantages: Captures non-linear relationships between features and produces more accurate imputations.
 - Disadvantages: Can overfit if not carefully validated and is computationally demanding.

In summary, handling missing data is an essential step in the feature engineering process. The choice of imputation technique depends on the nature of the data, the amount of missing data, and the computational resources available. Each technique has its strengths and weaknesses, so understanding the characteristics of the dataset and the underlying patterns of missingness is crucial for making the right decision.